# Independent sampling runs | AIME 2026

Two generation seeds · 256 responses per question · 30 questions per model

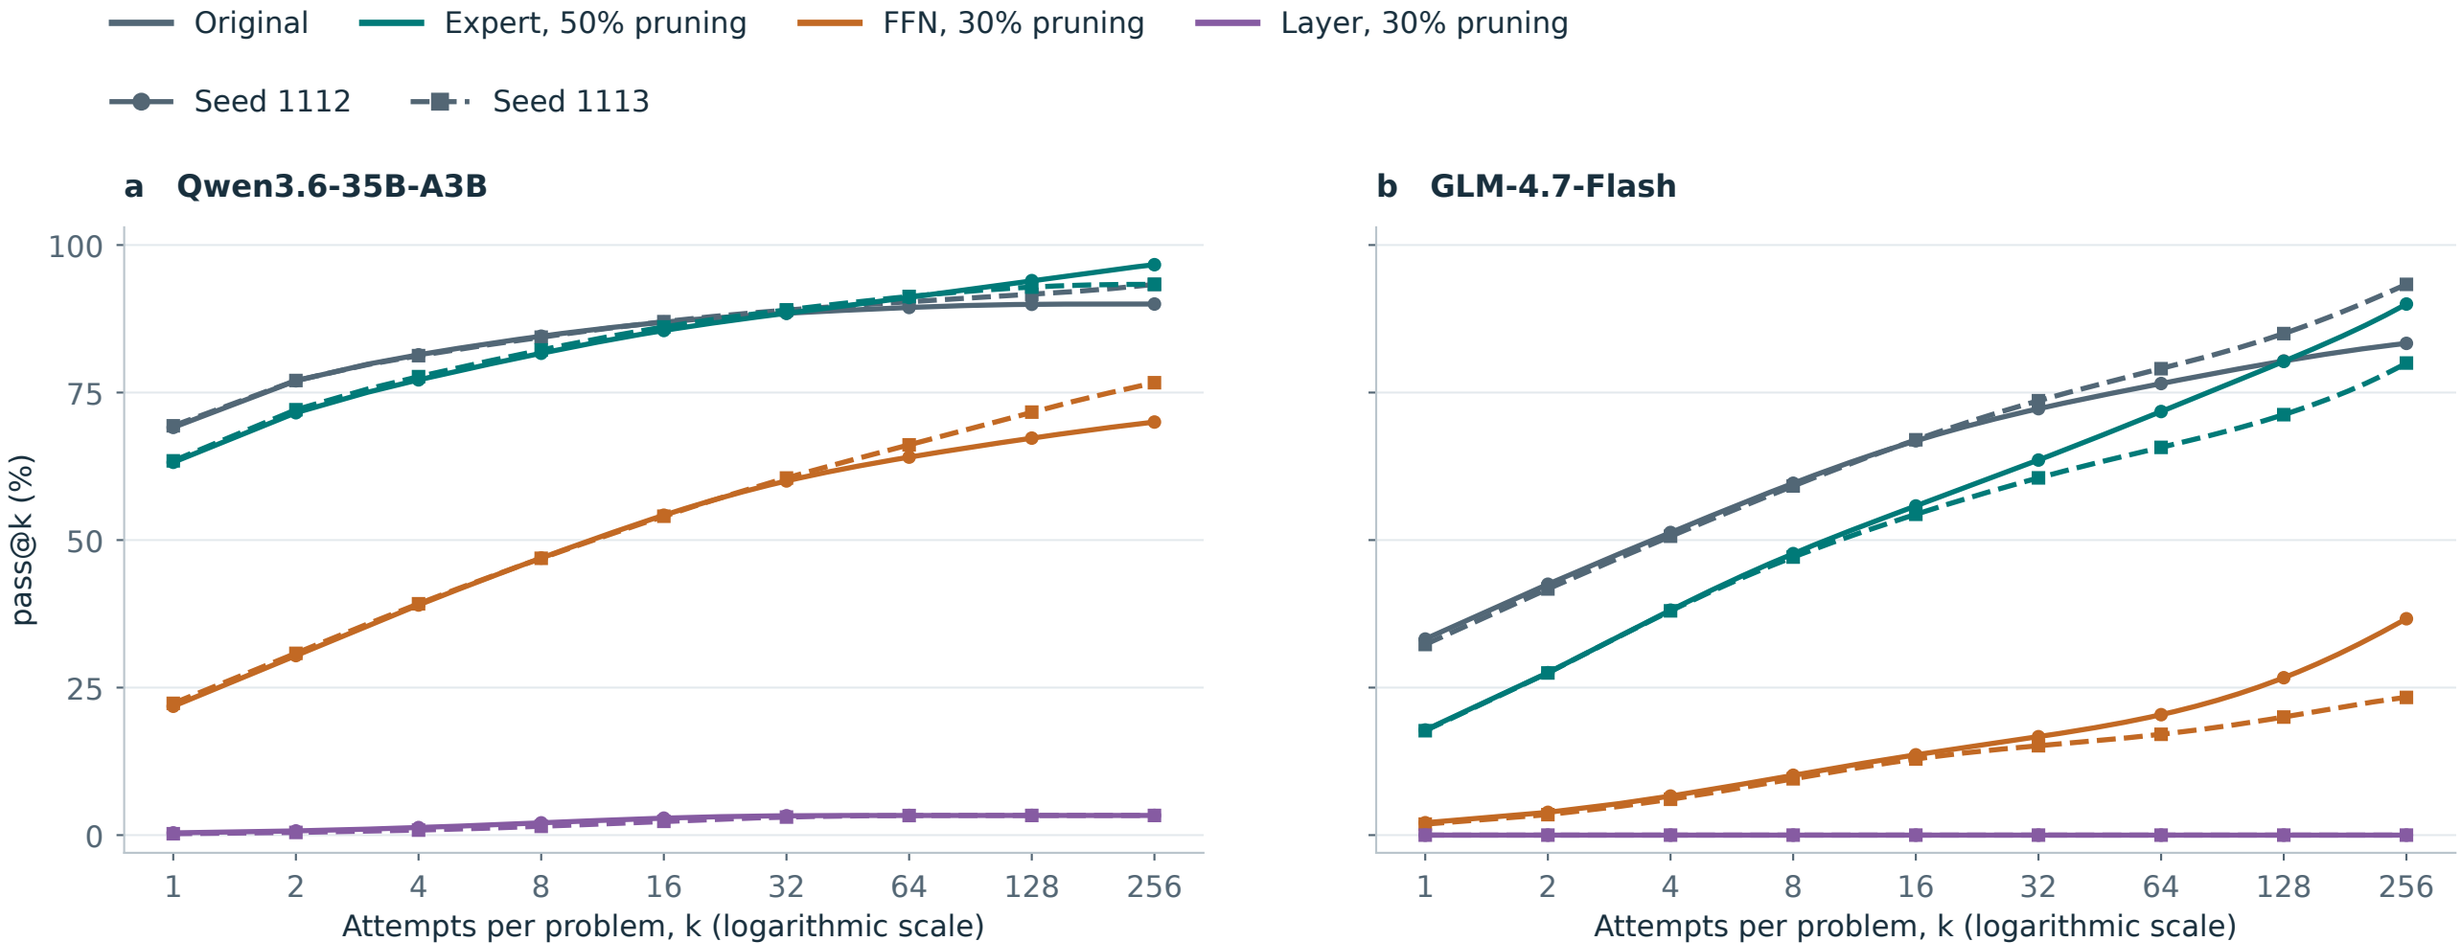


Temperature 0.7 · top-p 0.95 · decoding top-k 20 · repetition penalty 1.0 · 8,192-token output limit.

Curves show separate generation runs. At k = 256, pass@k is the fraction of problems solved at least once.