Figure 1 · Coverage at a fixed candidate budget

GLM-4.7-Flash | Original and 10%, 30%, 50% pruning

Quality: 8,192-token response limit; mathematics k = 2,048, code k = 512. Lower resource use is to the left.
FLOPs and timing: separate 512-input / 128-output workload, one H20, one request at a time. Timing: median of five batches.
Layer 30% has resource measurements but no corresponding quality result; its quality point is omitted.