Figure 1 · Reasoning coverage

GLM-4.7-Flash

Original · Expert · FFN · Layer · 10% pruning · 30% pruning · 50% pruning

Total parameters | Active parameters | Estimated FLOPs | Decode time

AIME 2026 pass@2,048 (%)

Minerva Math pass@2,048 (%)

HumanEval+ pass@512 (%)

LiveCodeBench-v6 pass@512 (%)

100
50
0

15 20 25 30
Billion parameters

2.0 2.5 3.0 3.5
Billion parameters/token

2.4 3.2 4.0
TFLOPs/request

3.0 4.5 6.0
ms/output token