Figure 1 · Coverage and resources

GLM-4.7-Flash | Original and 10%, 30%, 50% pruning

Quality: 8,192-token response limit. Each task uses the fixed sampling budget shown above.
Resources: one H20, 512 input / 128 output tokens, one request at a time. Latency is a five-batch median.
Layer 30% has resource measurements but no corresponding quality result; no point is drawn.