Figure 3 · Output length, success, and truncation

GLM-4.7-Flash | AIME 2026

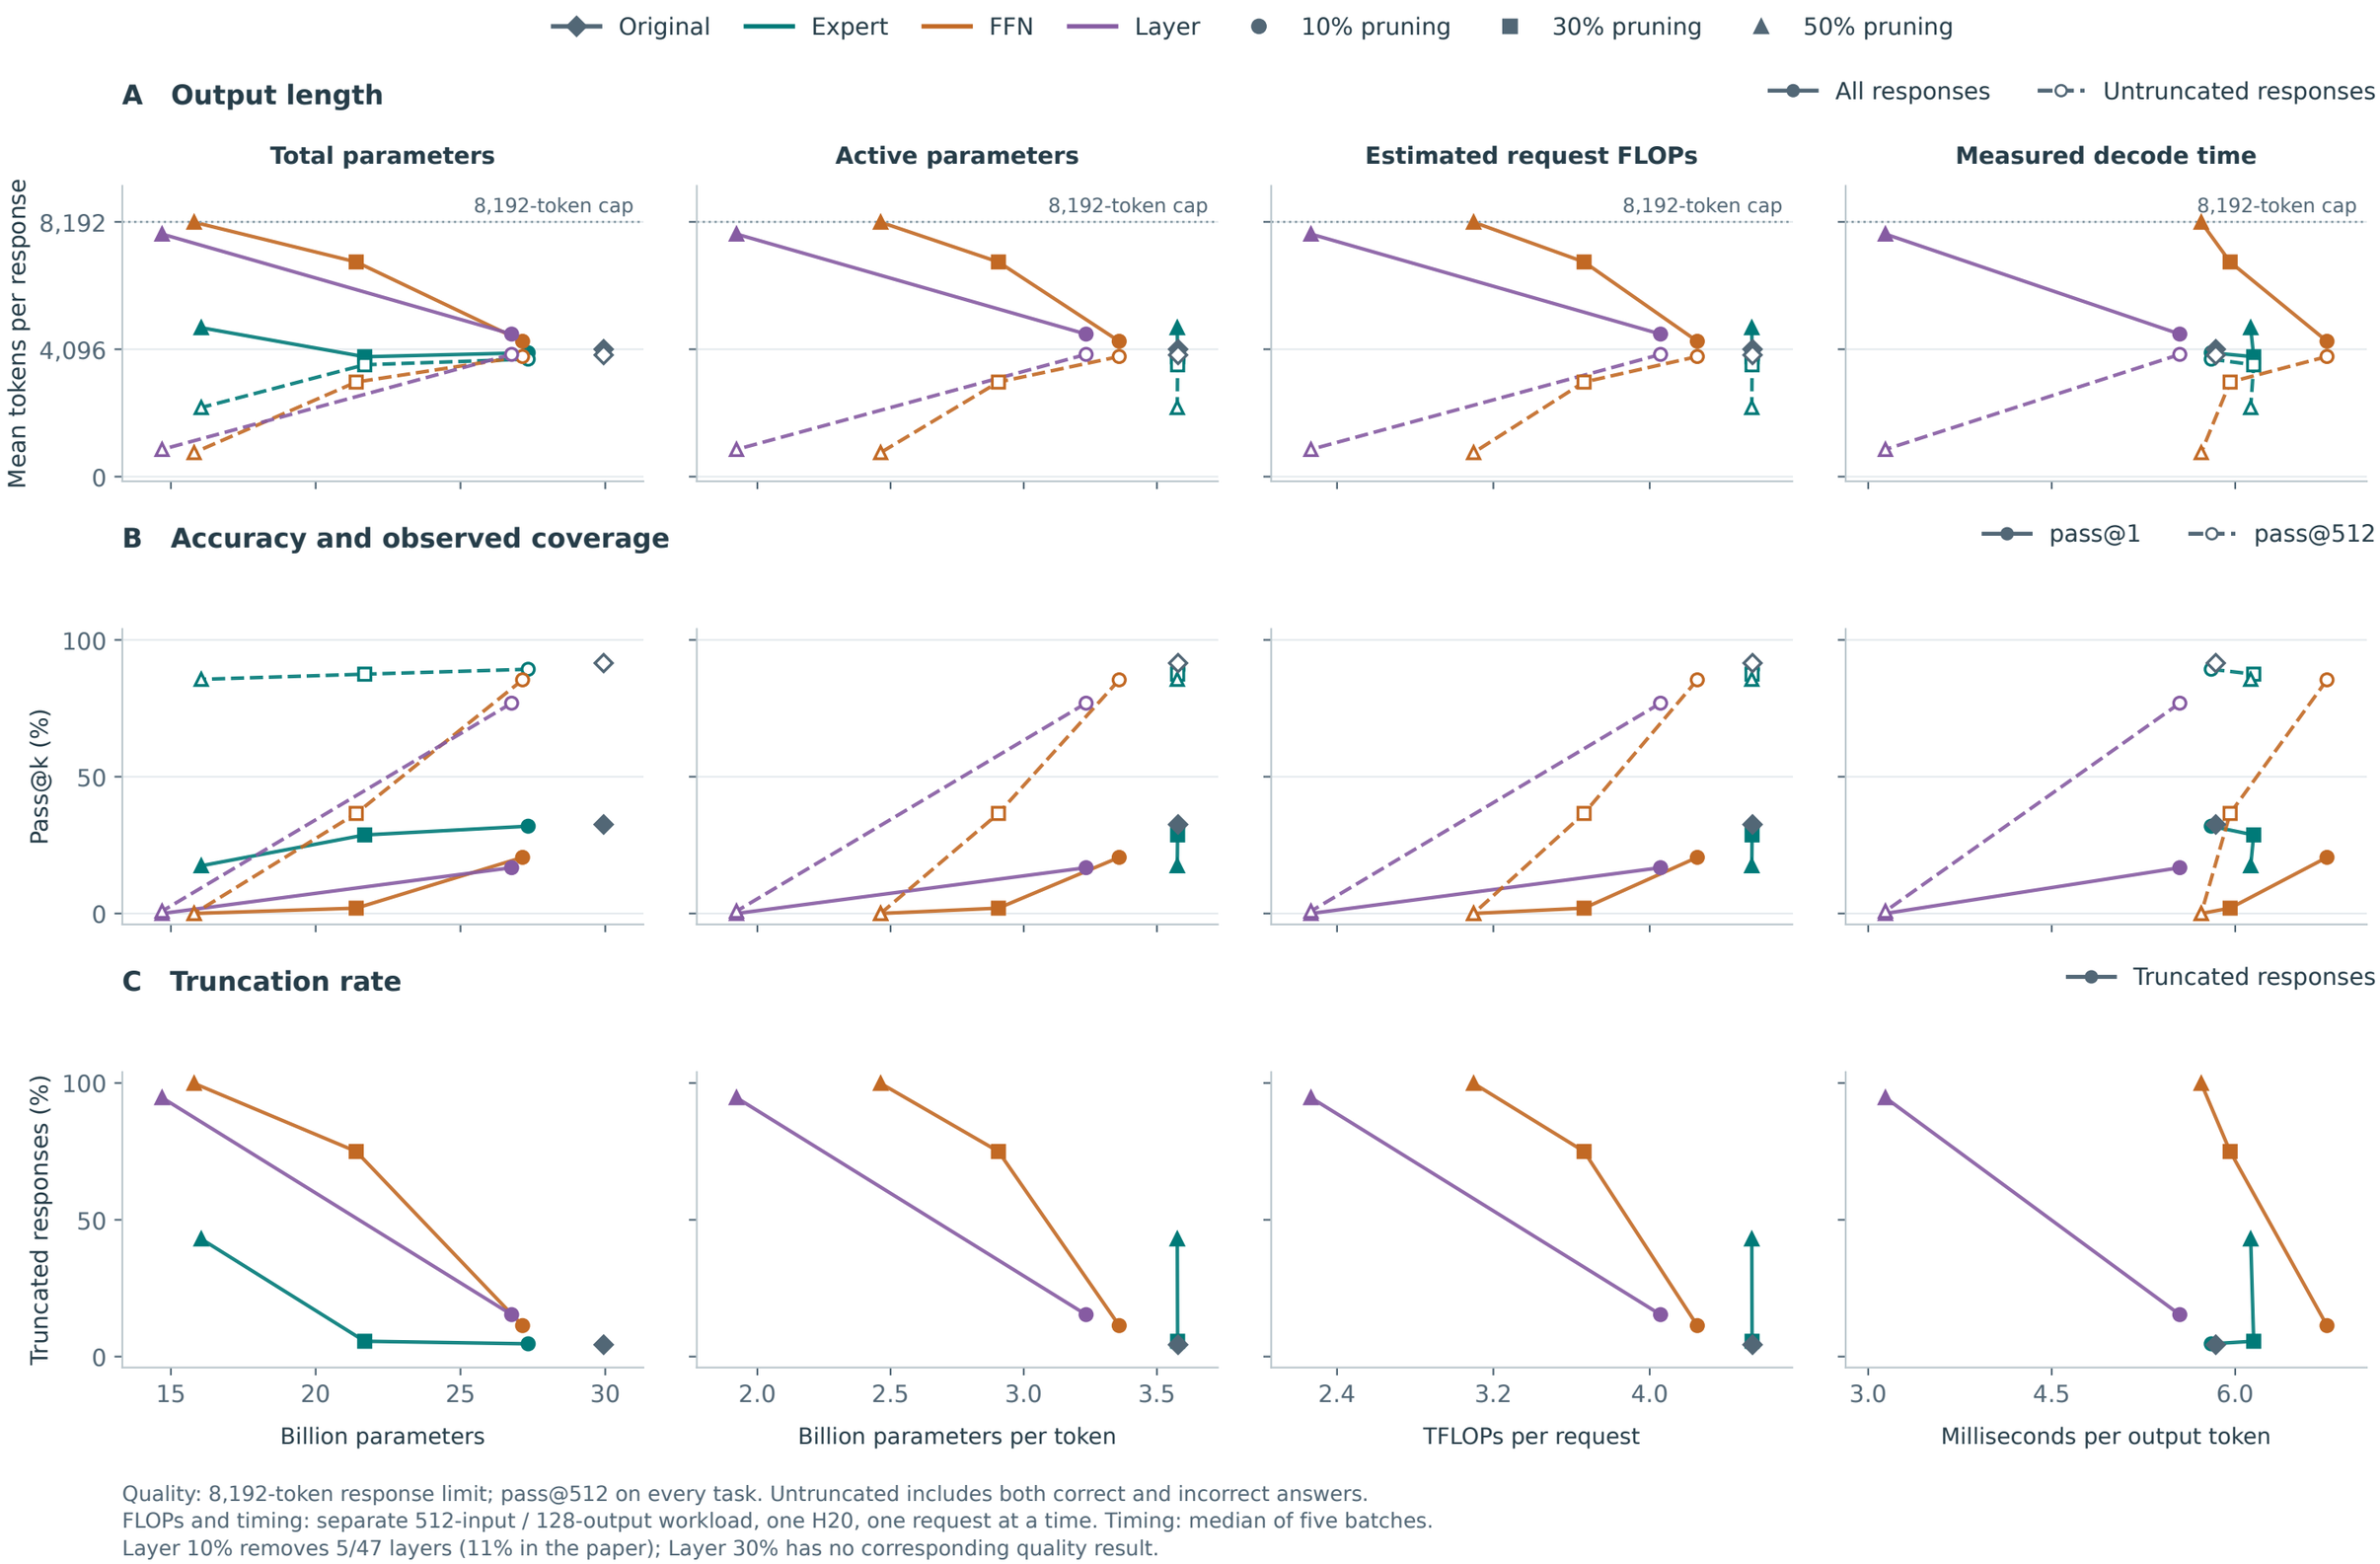

Quality: 8,192-token response limit; pass@512 on every task. Untruncated includes both correct and incorrect answers.
FLOPs and timing: separate 512-input / 128-output workload, one H20, one request at a time. Timing: median of five batches.
Layer 10% removes 5/47 layers (11% in the paper); Layer 30% has no corresponding quality result.