GLM-4.7-Flash · HumanEval+

Figure 3

Original · Expert · FFN · Layer

10% pruning · 30% pruning · 50% pruning

A Output length

All responses · Untruncated

Total parameters | Active parameters | Estimated FLOPs | Decode time

Mean output tokens: 8,192 · 4,096 · 0

B Accuracy and coverage

pass@1 · pass@512

Pass@k (%): 100 · 50 · 0

C Truncation rate

Truncated responses (%): 100 · 50 · 0

Billion parameters: 15 · 20 · 25 · 30

Billion parameters/token: 2.0 · 2.5 · 3.0 · 3.5

TFLOPs/request: 2.4 · 3.2 · 4.0

ms/output token: 3.0 · 4.5 · 6.0