Figure 3 · Output length, success, and truncation

Qwen3.6-35B-A3B | AIME 2026

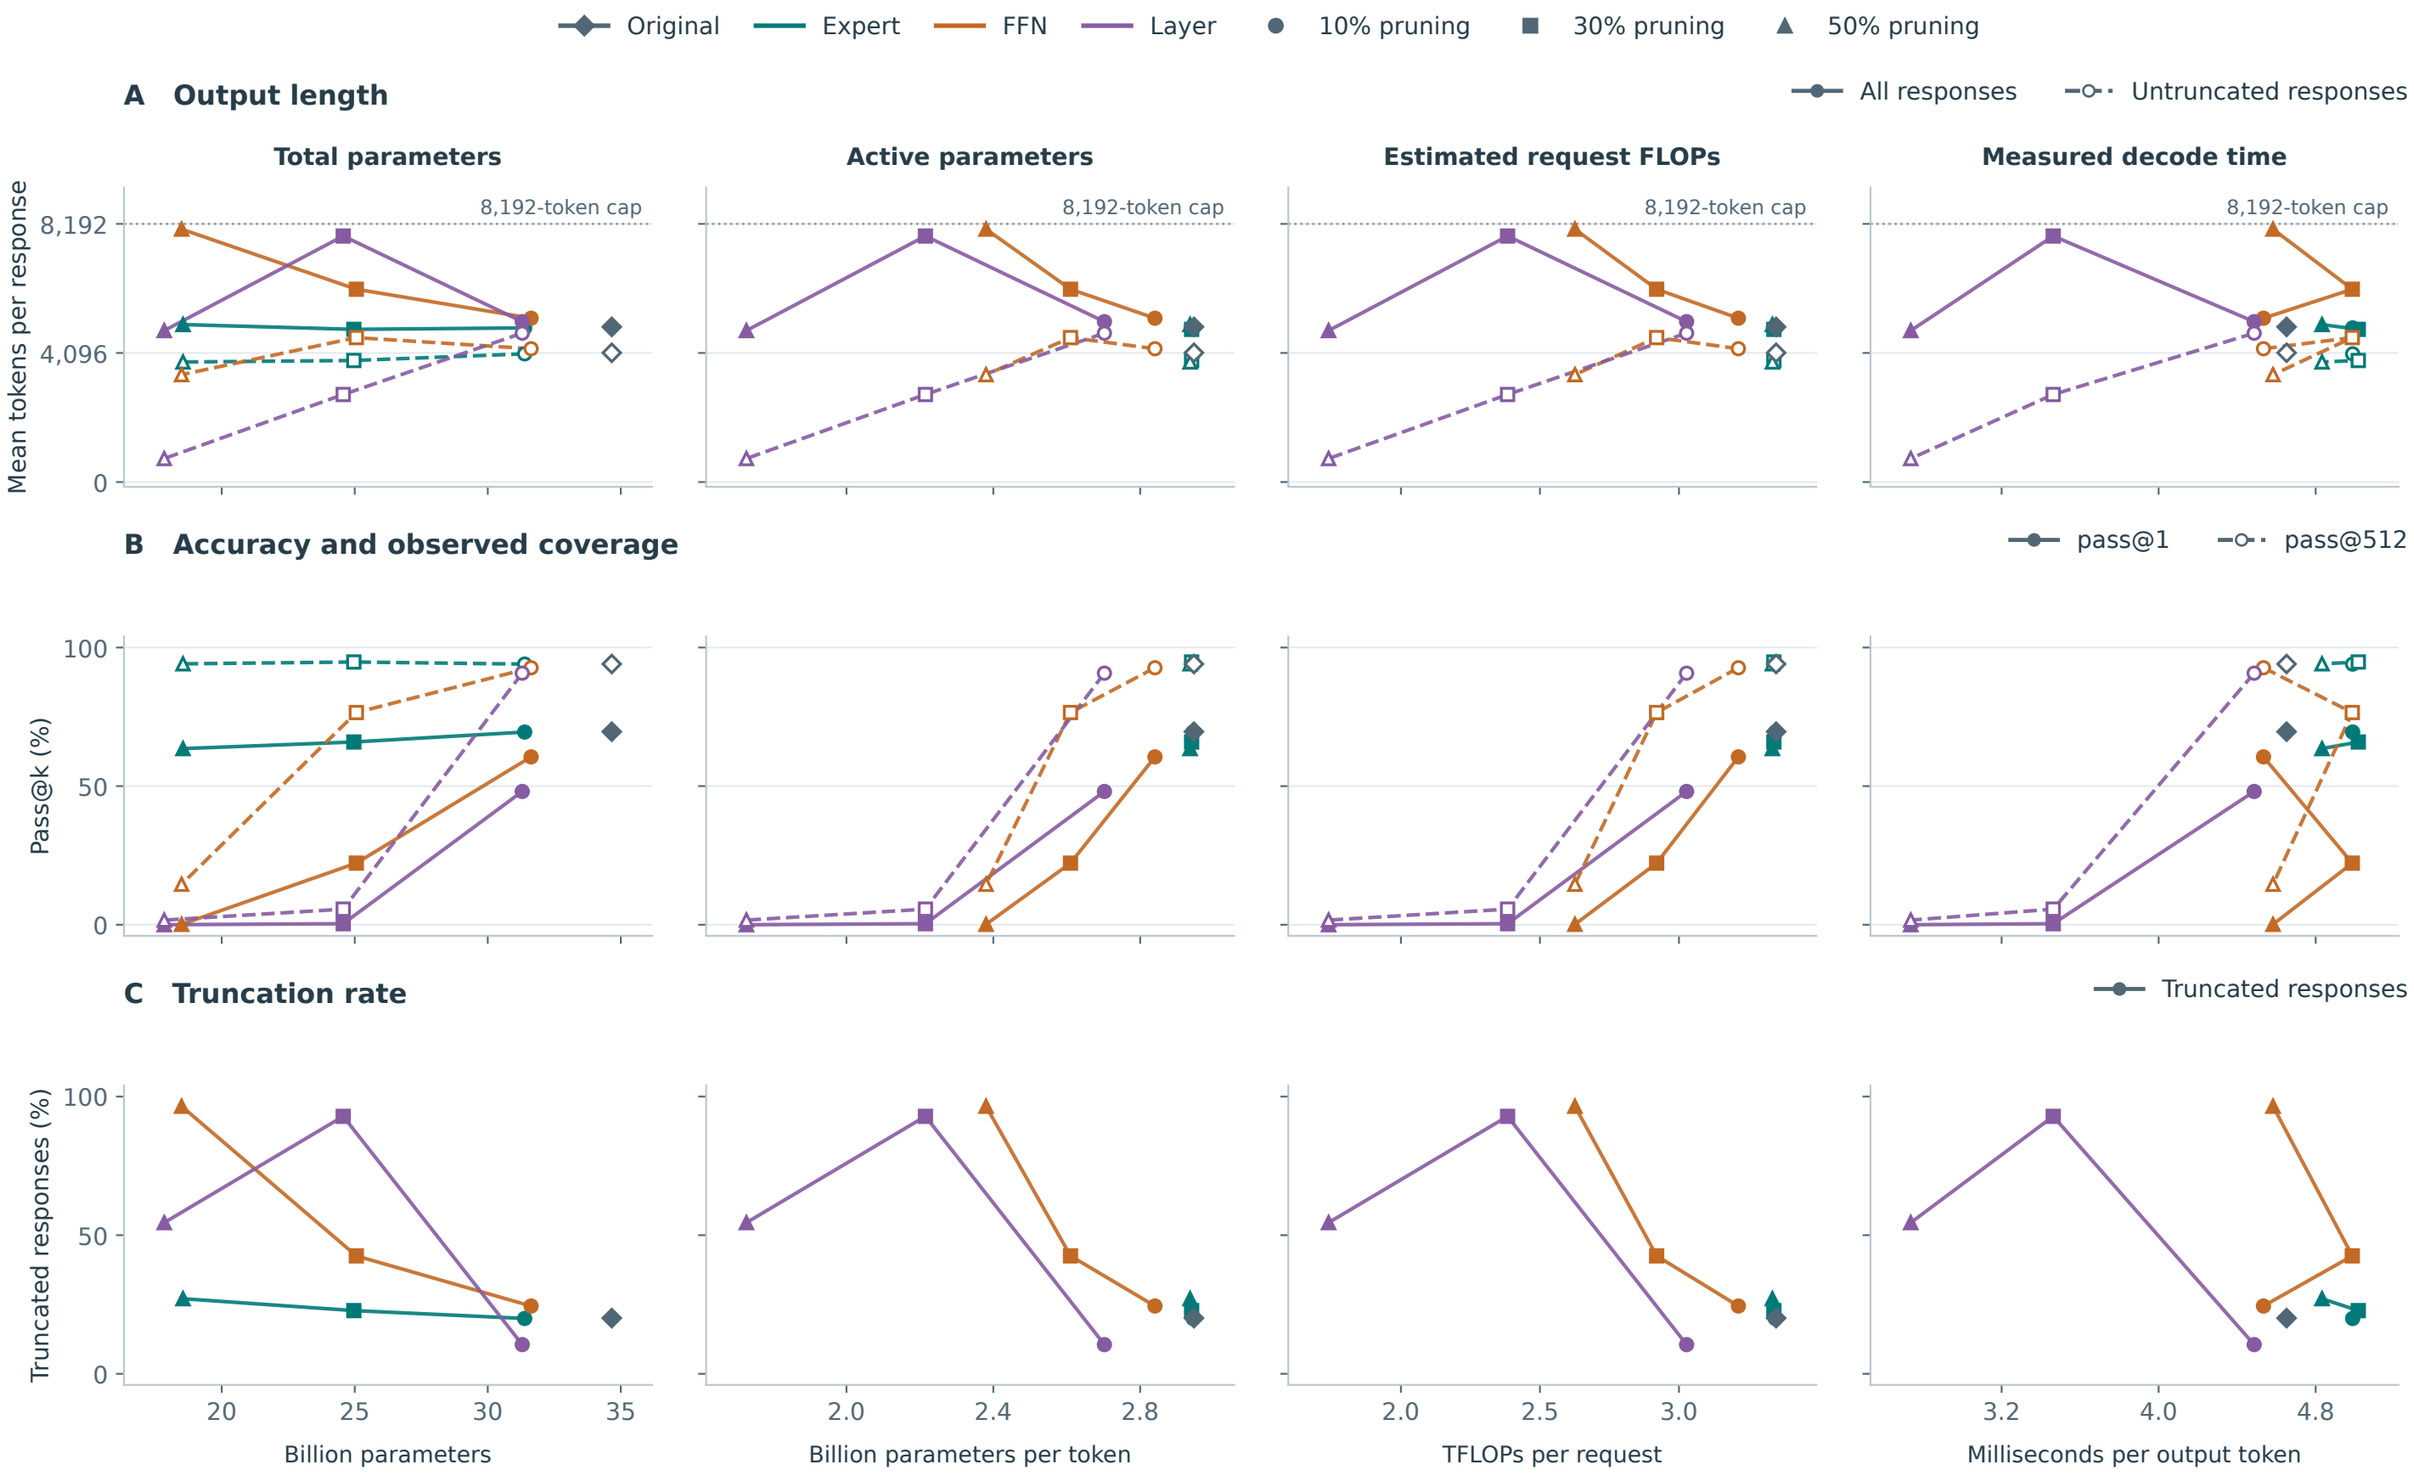

Quality: 8,192-token response limit; pass@512 on every task. Untruncated includes both correct and incorrect answers.
FLOPs and timing: separate 512-input / 128-output workload, one H20, one request at a time. Timing: median of five batches.