Figure 3 · Length, quality, and resources

Qwen3.6-35B-A3B | AIME 2026

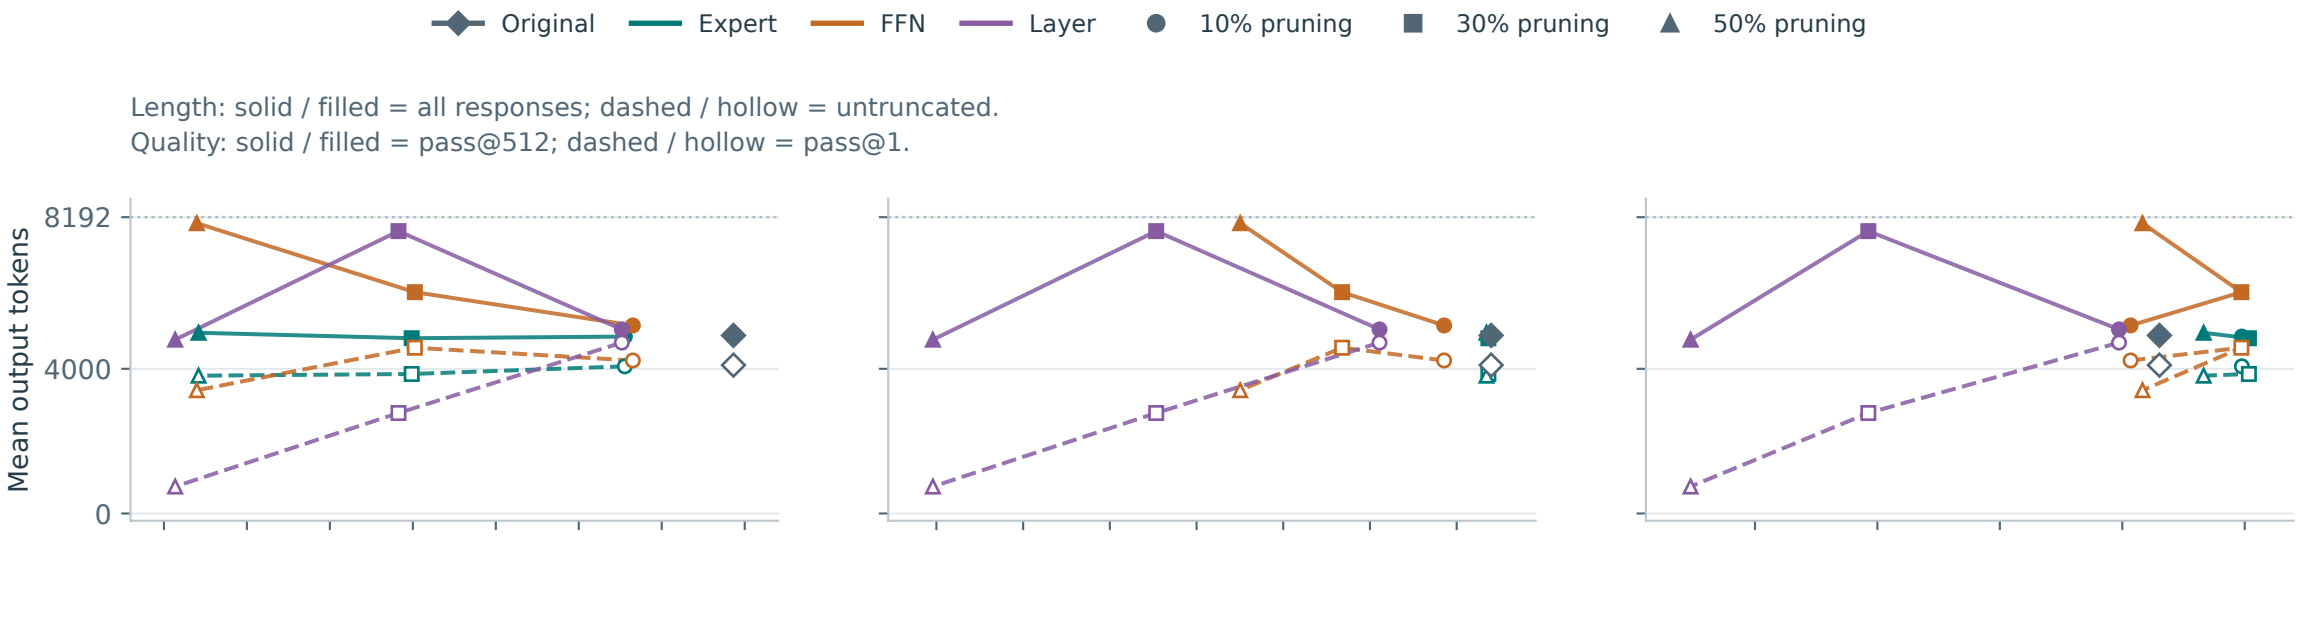

Quality: 8,192-token response limit; pass@512 on every task. Untruncated includes incorrect answers.
Resources: one H20, 512 input / 128 output tokens, one request at a time. Latency is a five-batch median.
Lines connect measured configurations in pruning-ratio order; no values are interpolated for missing configurations.