Qwen3.6-35B-A3B · Minerva Math

Figure 3

Original · Expert · FFN · Layer

10% pruning · 30% pruning · 50% pruning

A Output length

All responses · Untruncated

Total parameters | Active parameters | Estimated FLOPs | Decode time

Mean output tokens

8,192

4,096

0

B Accuracy and coverage

pass@1 · pass@512

Pass@k (%)

100

50

0

C Truncation rate

Truncated responses (%)

100

50

0

20 25 30 35

Billion parameters

2.0 2.4 2.8

Billion parameters/token

2.0 2.5 3.0

TFLOPs/request

3.2 4.0 4.8

ms/output token