Figure 3 · Output length, success, and truncation

Qwen3.6-35B-A3B | Minerva Math

Original · Expert · FFN · Layer · 10% pruning · 30% pruning · 50% pruning

A Output length

All responses · Untruncated responses

Total parameters | Active parameters | Estimated request FLOPs | Measured decode time

Mean tokens per response

8,192-token cap

8,192 · 4,096 · 0

B Accuracy and observed coverage

pass@1 · pass@512

Pass@k (%)

100 · 50 · 0

C Truncation rate

Truncated responses

Truncated responses (%)

100 · 50 · 0

20 · 25 · 30 · 35

Billion parameters

2.0 · 2.4 · 2.8

Billion parameters per token

2.0 · 2.5 · 3.0

TFLOPs per request

3.2 · 4.0 · 4.8

Milliseconds per output token

Quality: 8,192-token response limit; pass@512 on every task. Untruncated includes both correct and incorrect answers.
FLOPs and timing: separate 512-input / 128-output workload, one H20, one request at a time. Timing: median of five batches.