GLM-4.7-Flash | AIME 2026

Output length, success, and truncation across pruning ratios

A Mean output length

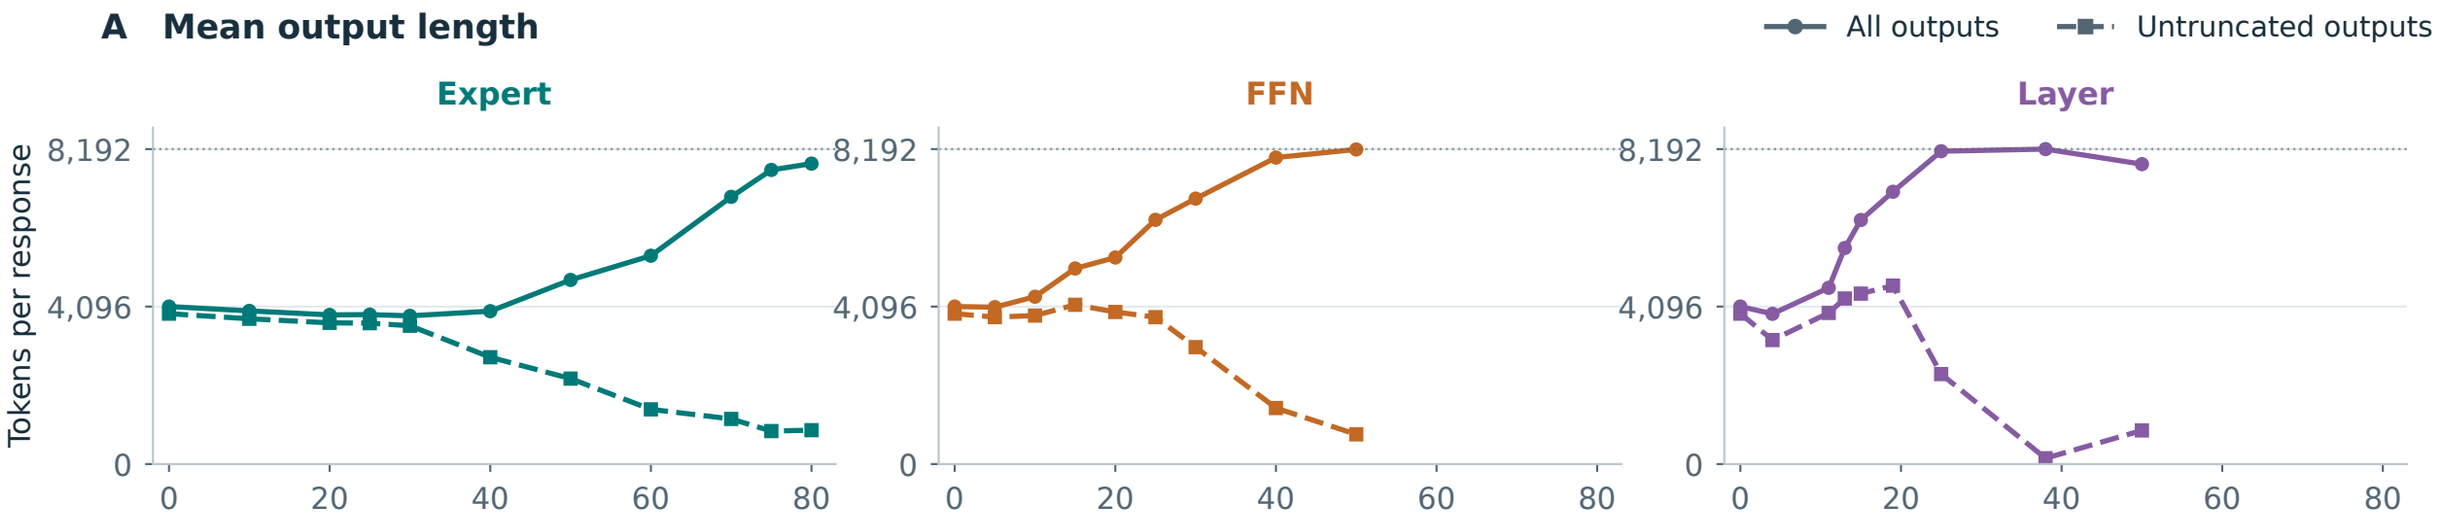

B Accuracy and observed coverage

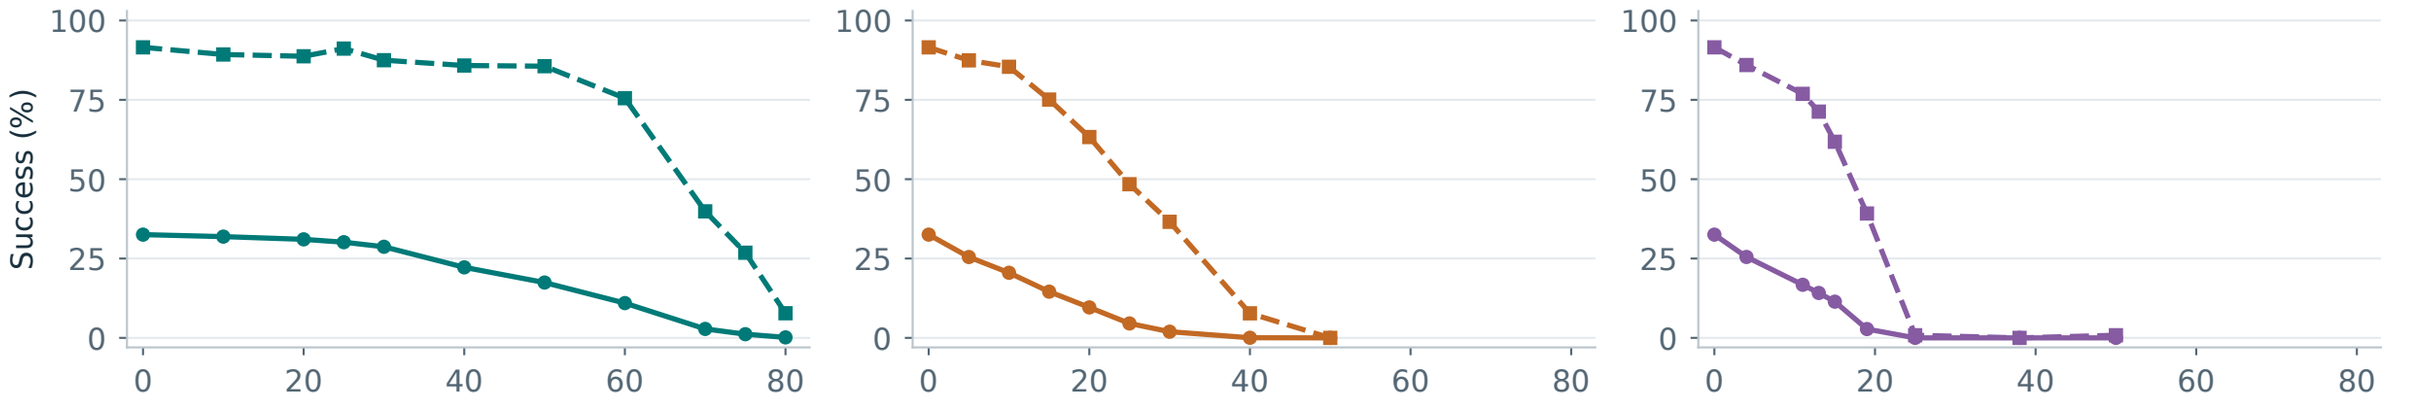

C Outputs reaching the token limit

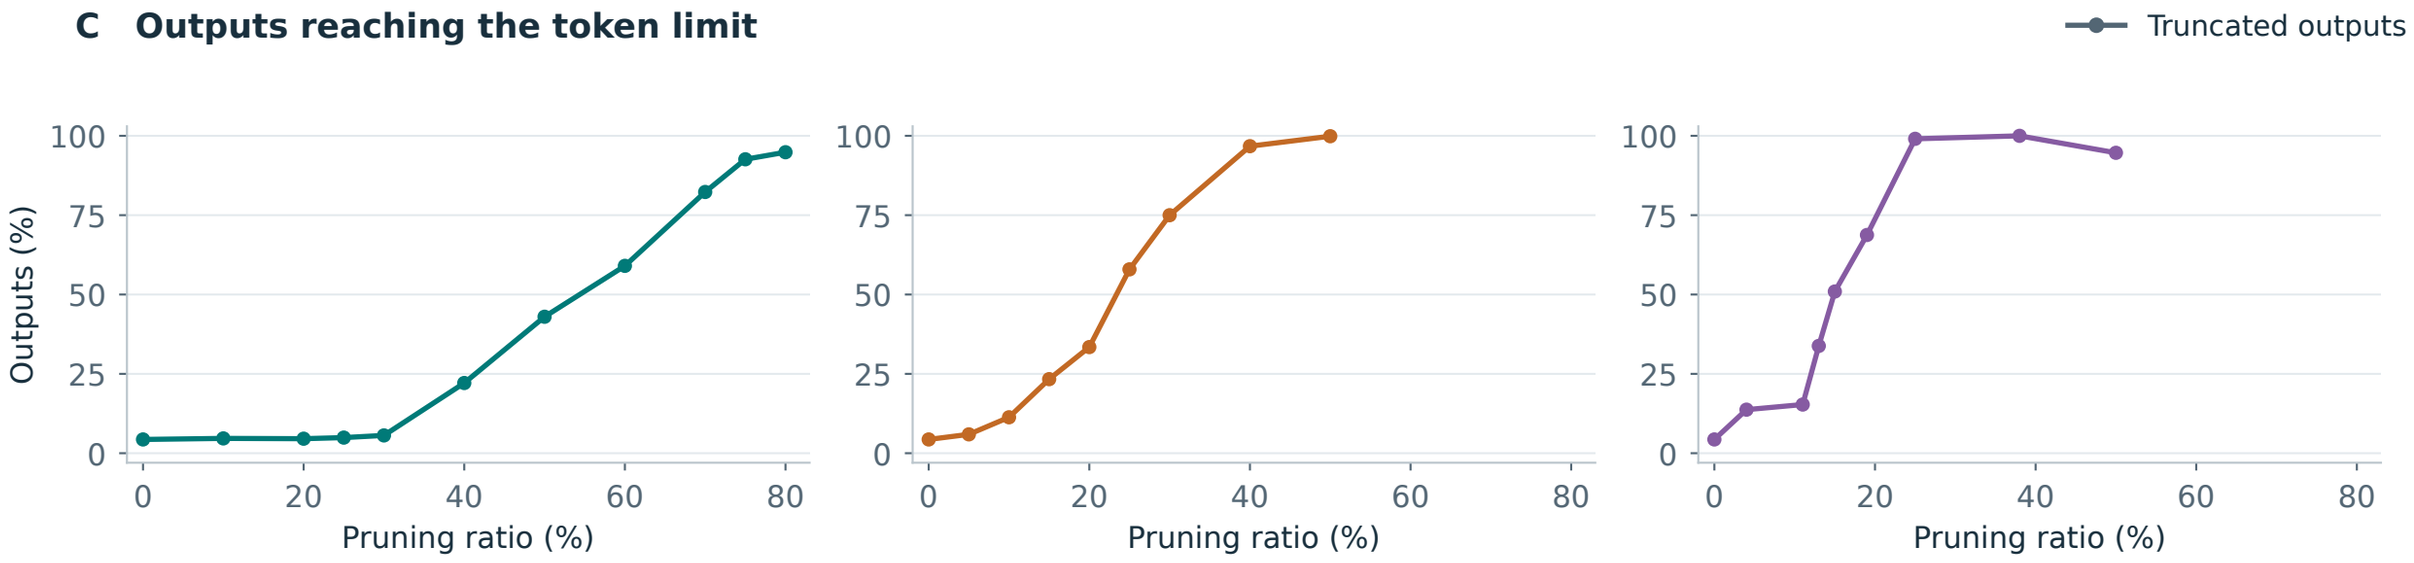

8,192-token output limit. Untruncated = below the limit, regardless of correctness.

Points retain the displayed pruning ratios; pass@512 uses the same number of attempts across all tasks.