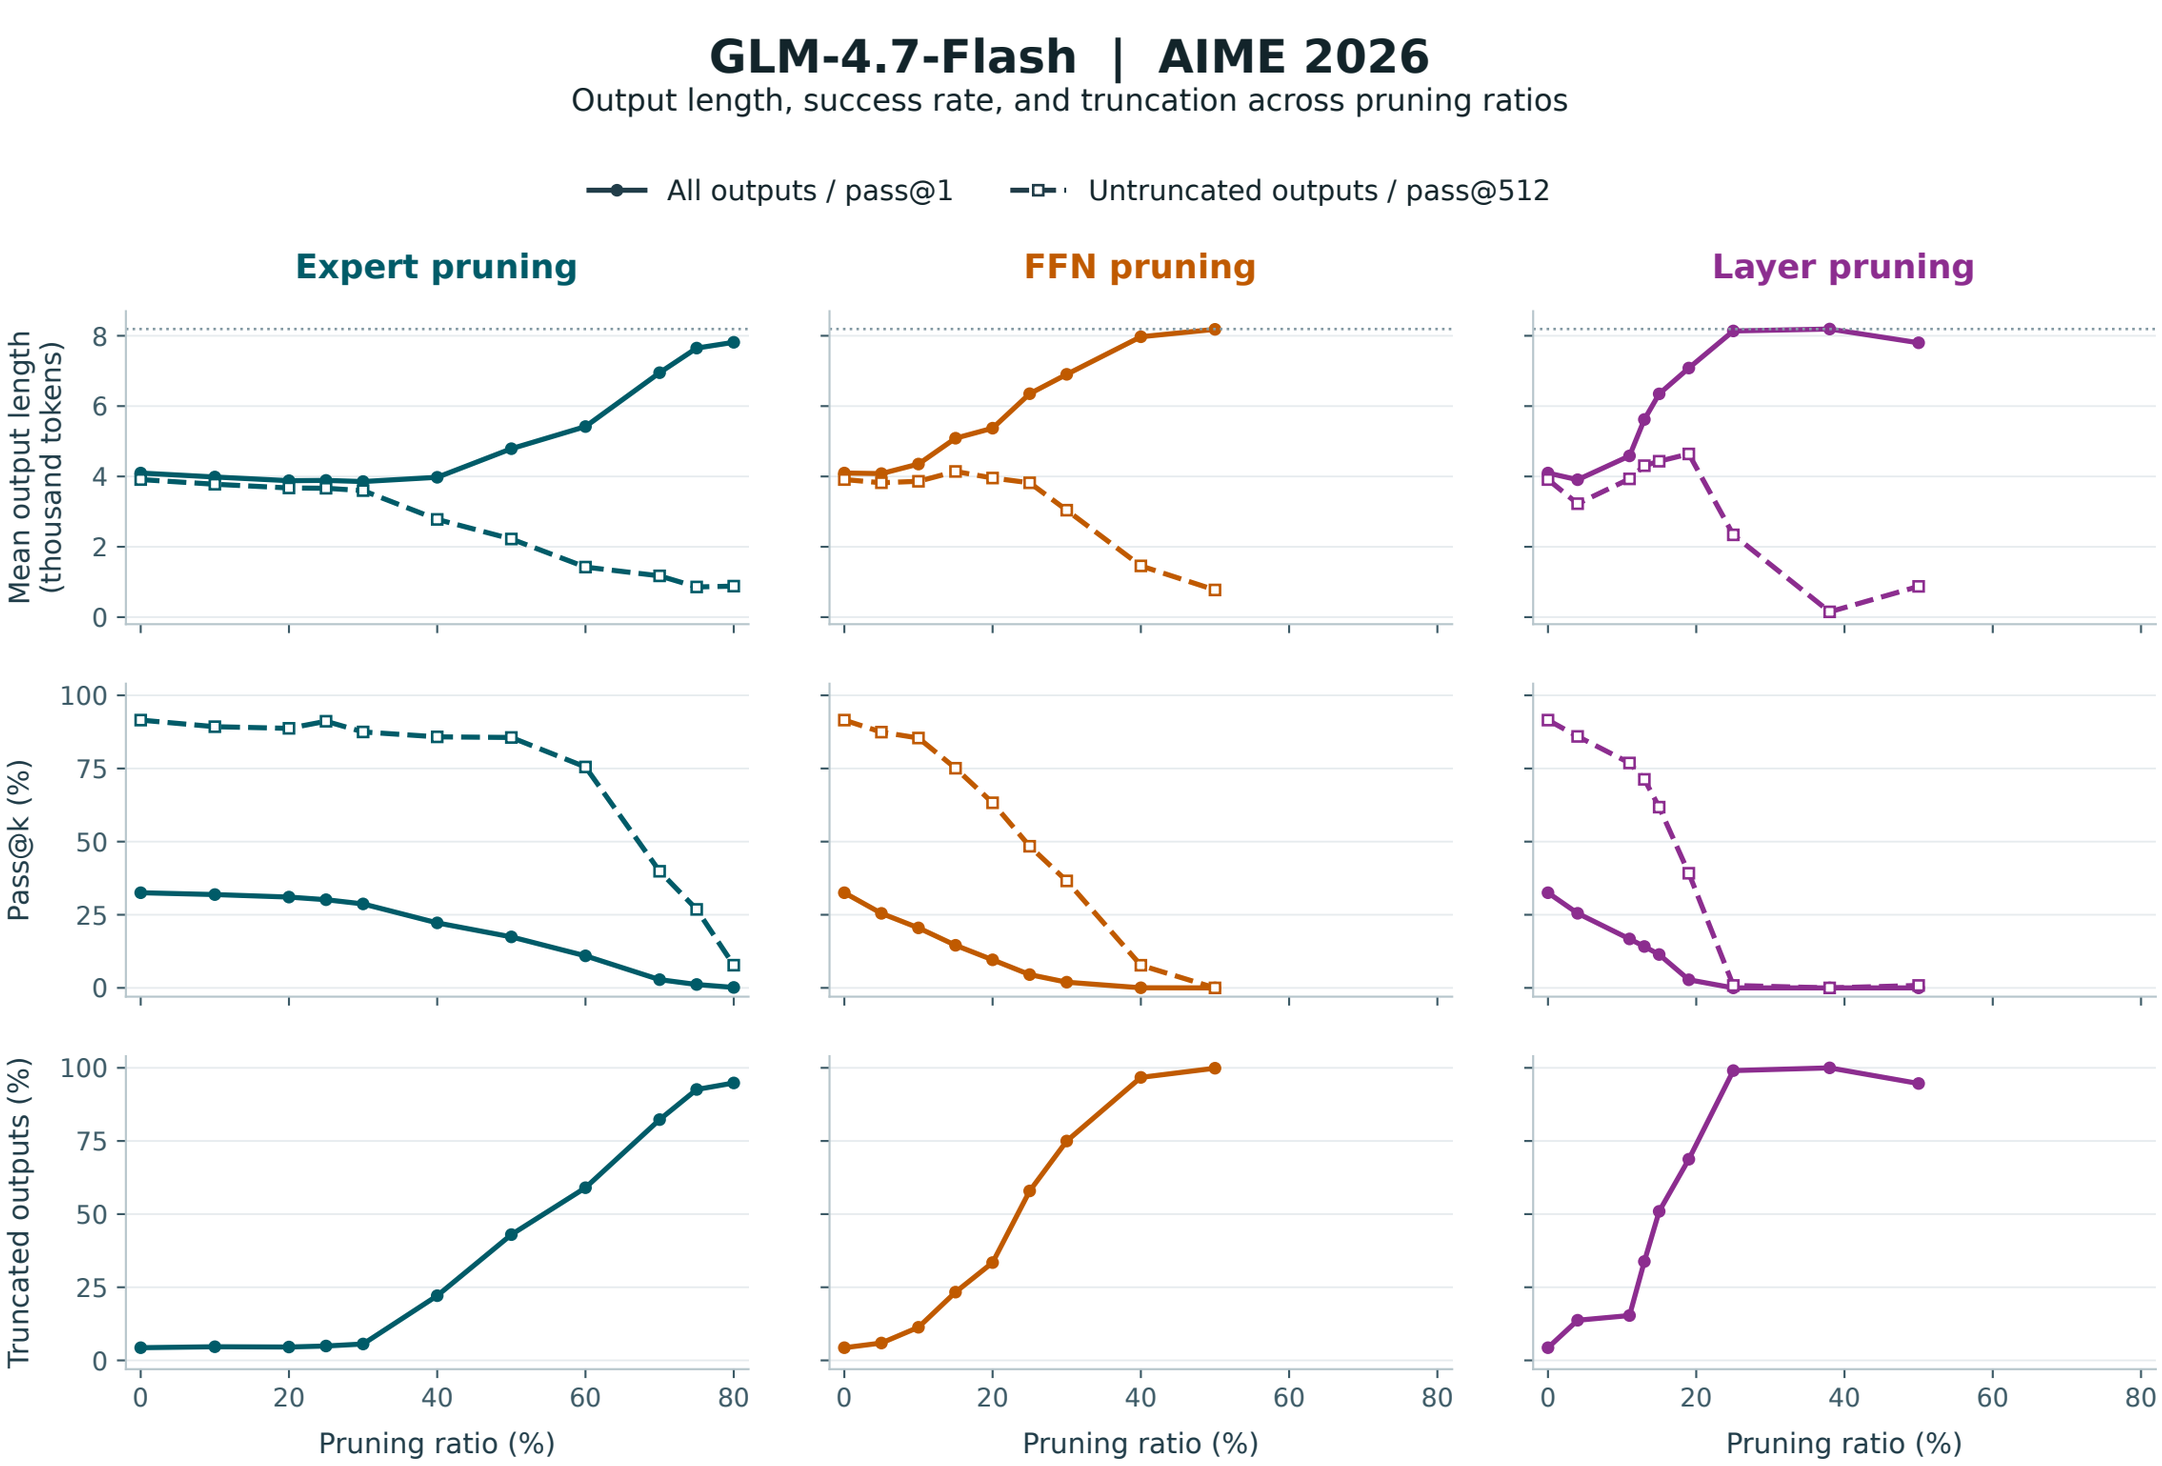
GLM-4.7-Flash | AIME 2026
Output length, success rate, and truncation across pruning ratios
All outputs / pass@1
Untruncated outputs / pass@512
Expert pruning
FFN pruning
Layer pruning
Mean output length (thousand tokens)
Pass@k (%)
Truncated outputs (%)
Pruning ratio (%)
Pruning ratio (%)
Pruning ratio (%)
8,192-token output limit. Untruncated means below the limit; it does not mean correct.