Figure 3 · Output length, success, and truncation

GLM-4.7-Flash | HumanEval+

Full pruning-ratio sweep | Separate scales for length, success, and truncation

Quality: 8,192-token response limit; pass@512 on every task. Untruncated includes both correct and incorrect answers.
Points retain the paper's displayed pruning ratios, including rounding for discrete structures.