Reasoning coverage and resource use

Original models and 10%, 30%, 50% pruning · Qwen and GLM

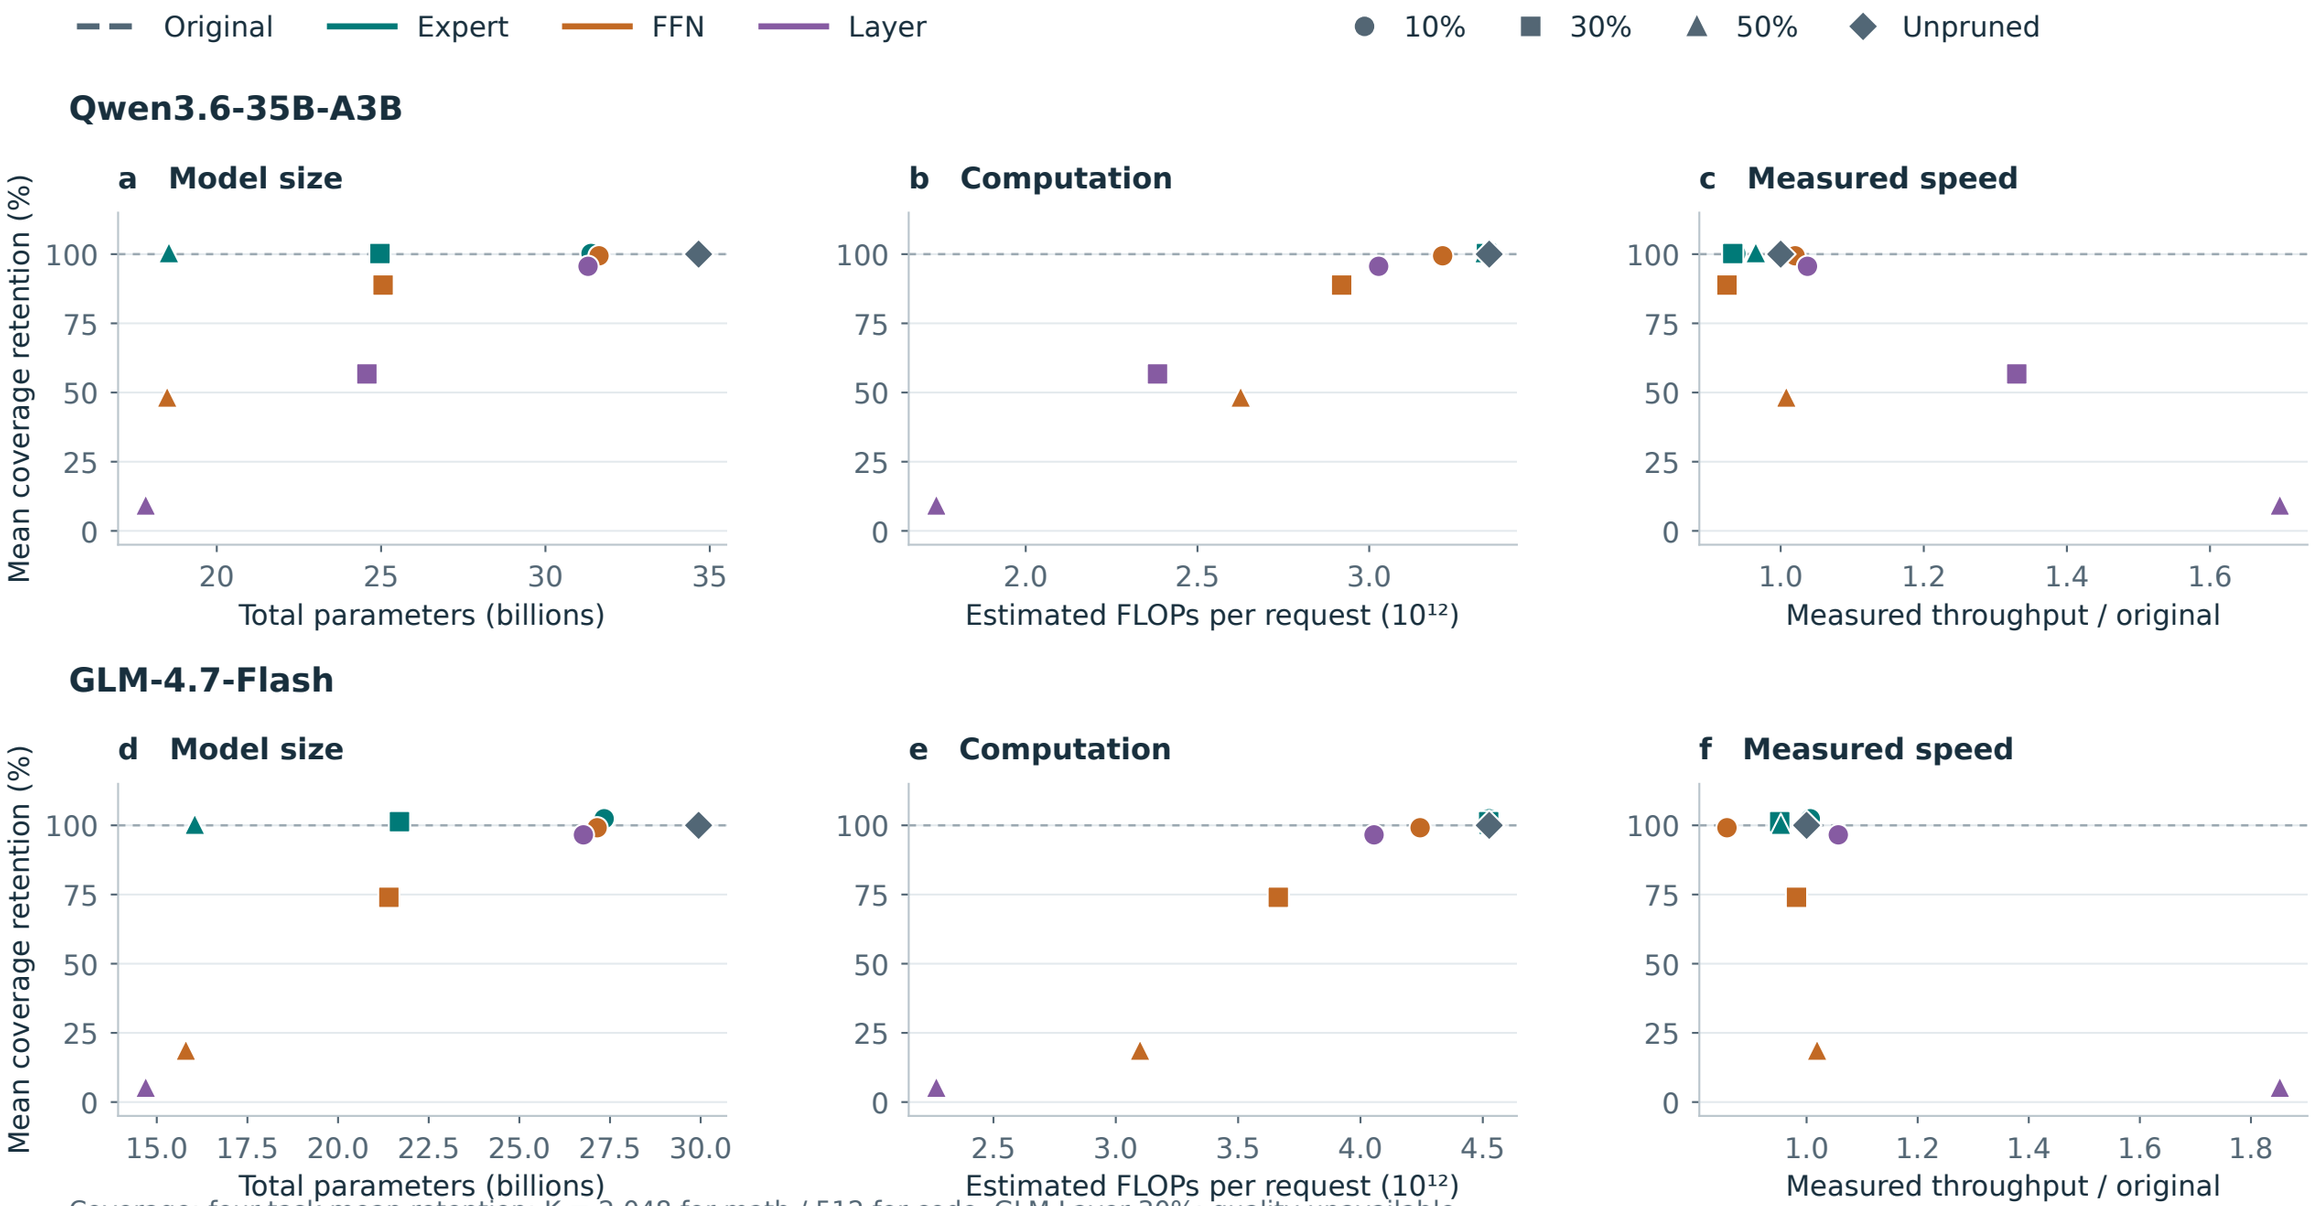

Coverage: four-task mean retention; K = 2,048 for math / 512 for code. GLM Layer 30%: quality unavailable.

Resource workload: 512 input / 128 output tokens, one request at a time; separate from reasoning generation.