Reasoning coverage and resource use

10%, 30%, and 50% pruning configurations

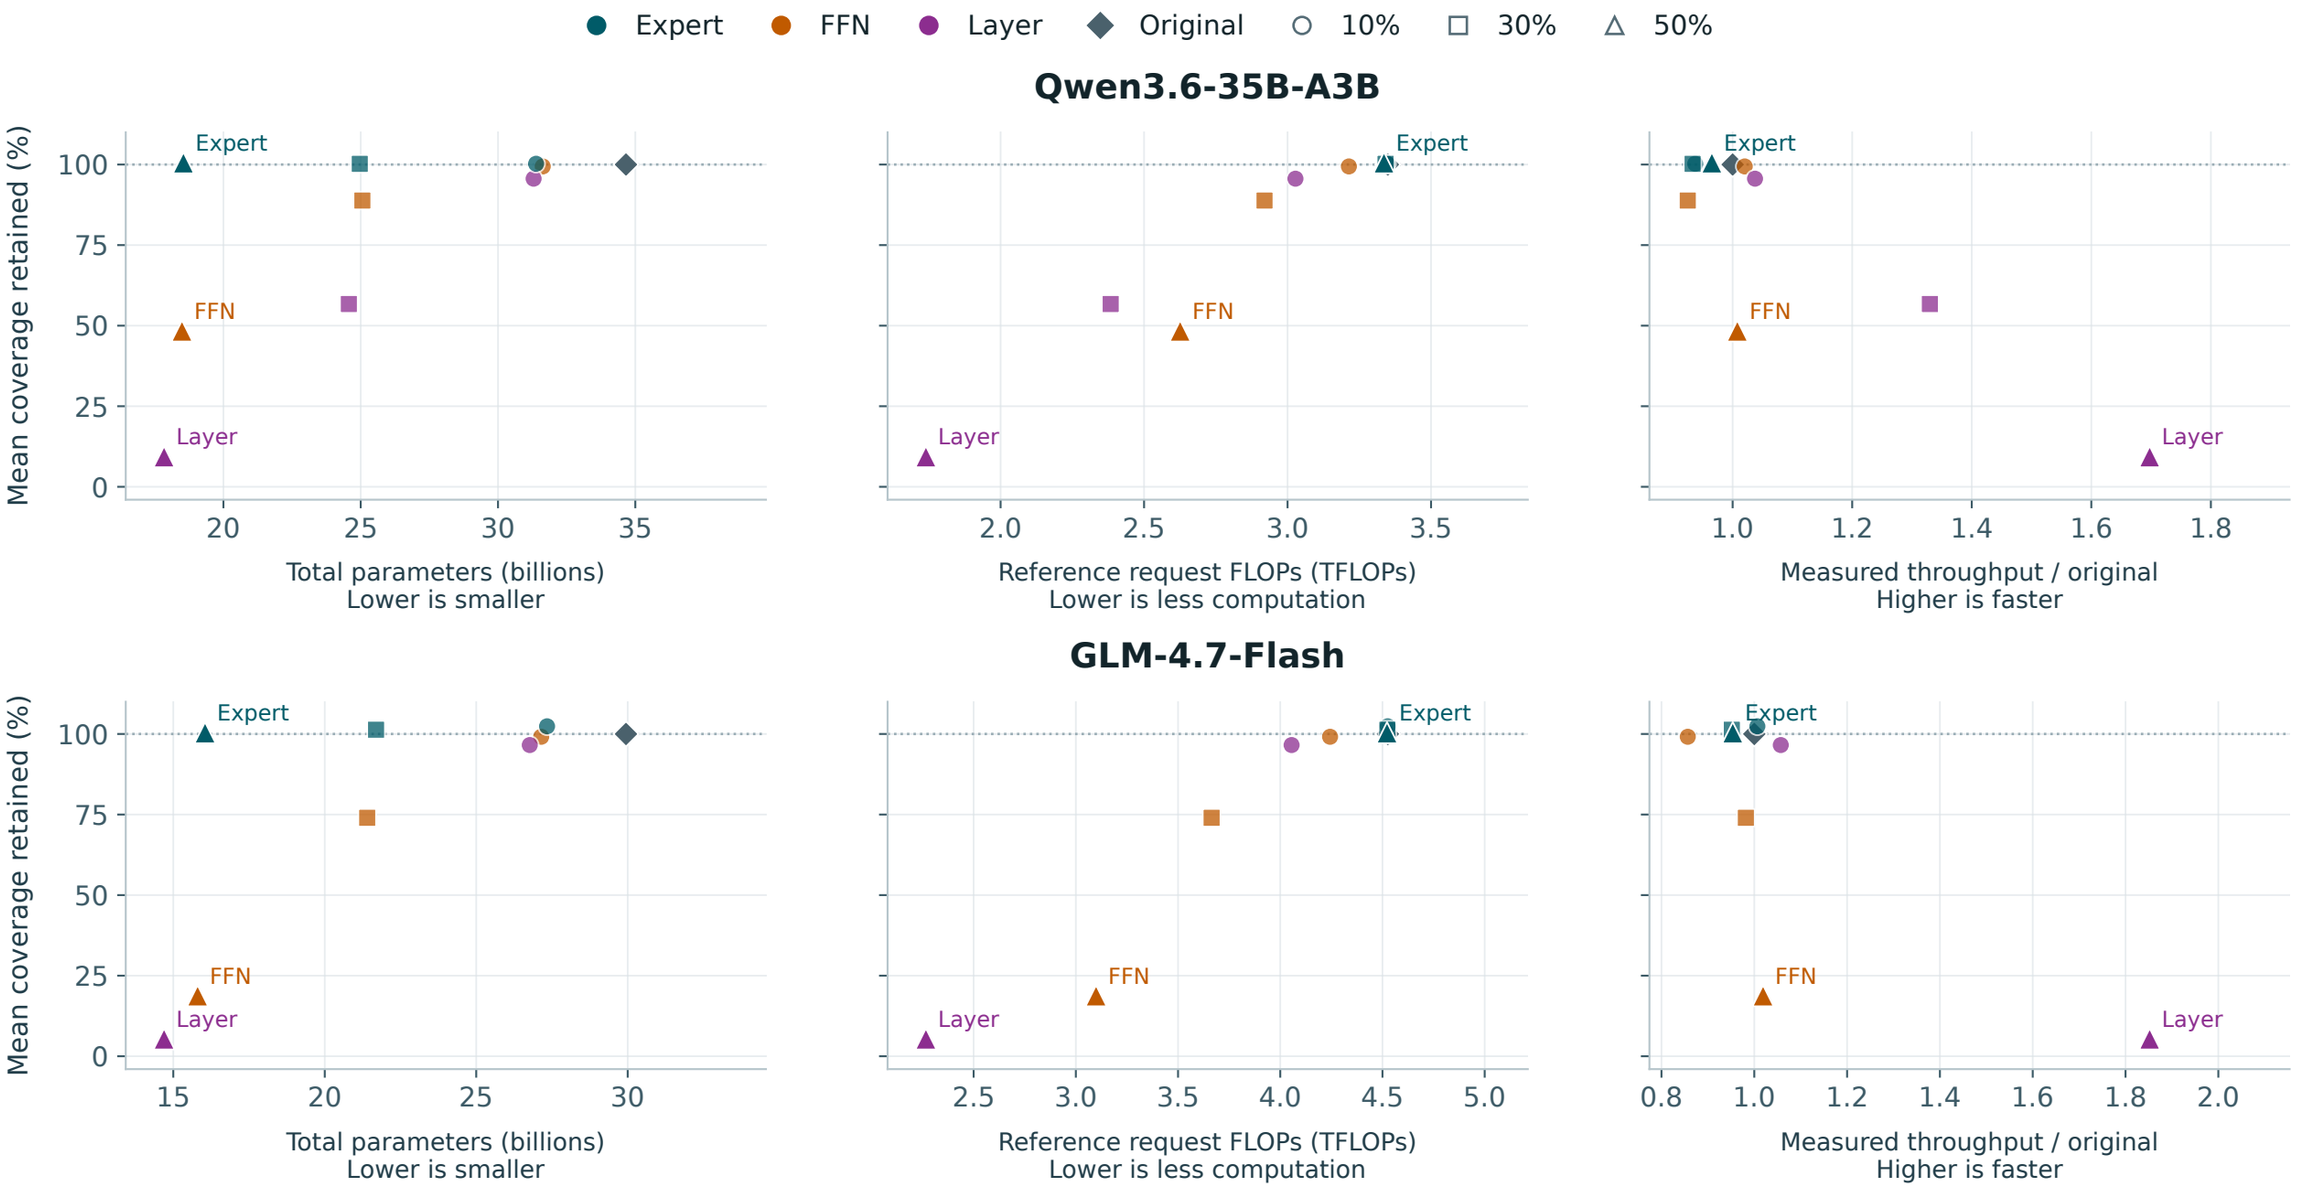

Coverage: mean of four taskwise ratios; math pass@2,048, code pass@512. Resource benchmark: 512 input + 128 output tokens, one request.

GLM Layer 30% has resource measurements only and is omitted here. Rebuilt 10% / 30% models reuse the historical pruning selections.