What each pruning structure saves

Original models and 10%, 30%, 50% pruning · Qwen and GLM

Original Expert FFN Layer

Qwen3.6-35B-A3B

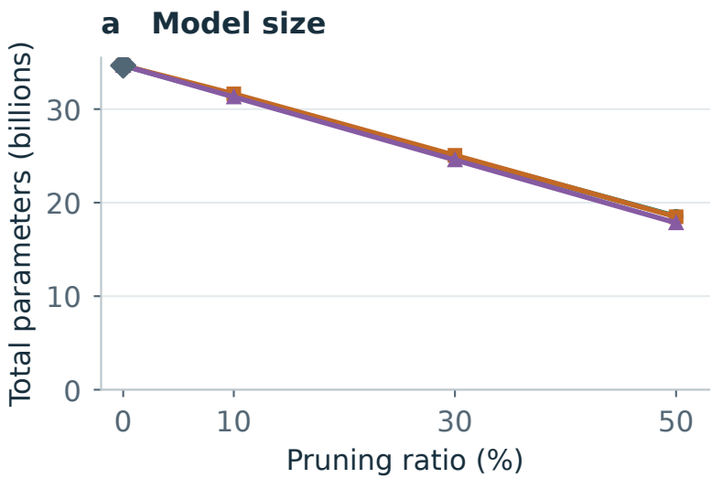

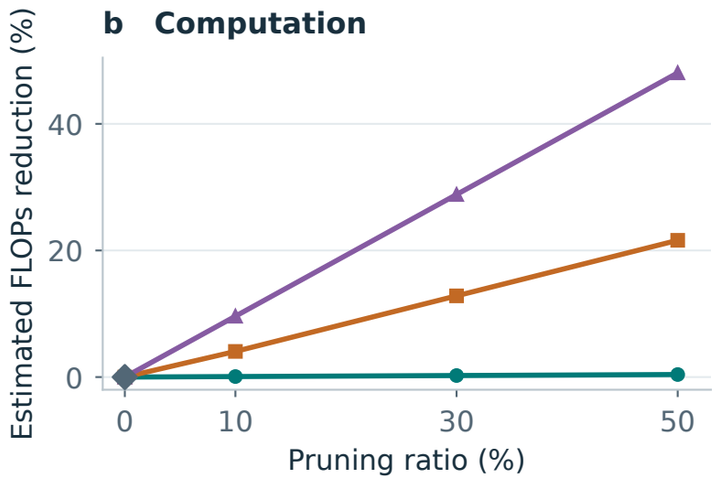

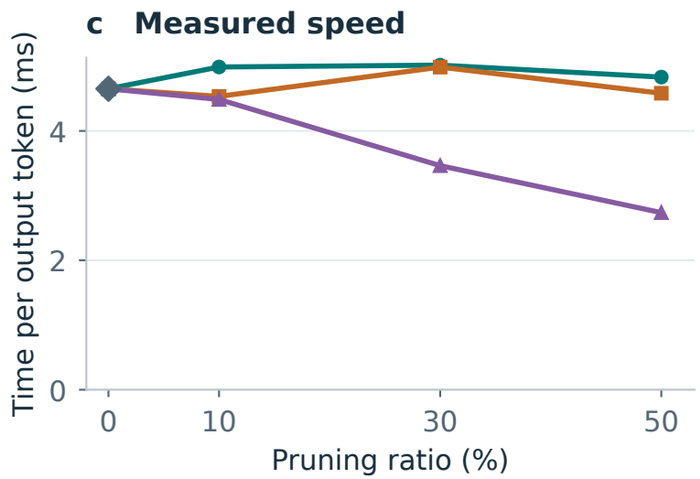

GLM-4.7-Flash

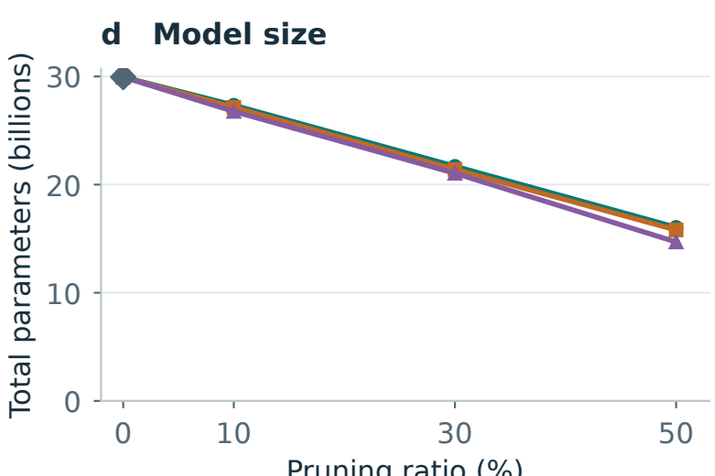

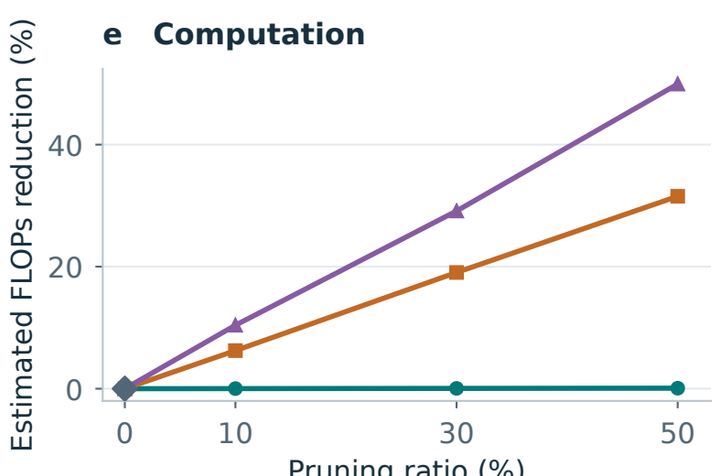

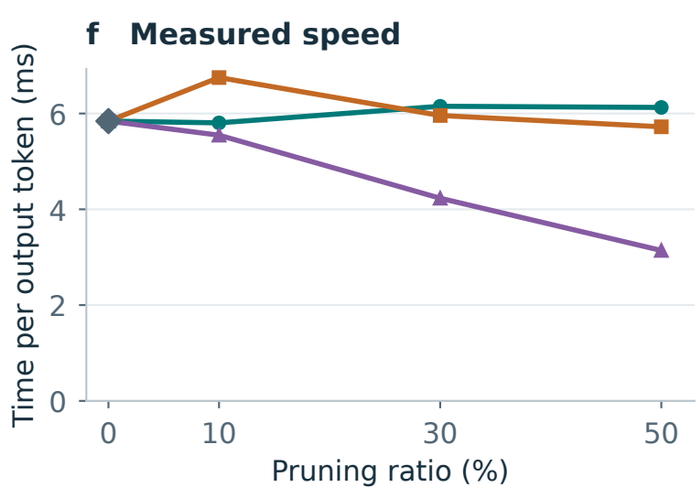

One H20 · BF16 · 512 input / 128 output tokens · one request at a time. Latency: median and IQR of five batches.
FLOPs are architecture-based counts. GLM Layer 10% / 30% remove 5 / 14 of 47 layers.