# Resource measurements across pruning ratios

H20 · BF16 · vLLM 0.25.1 · one GPU · one request at a time

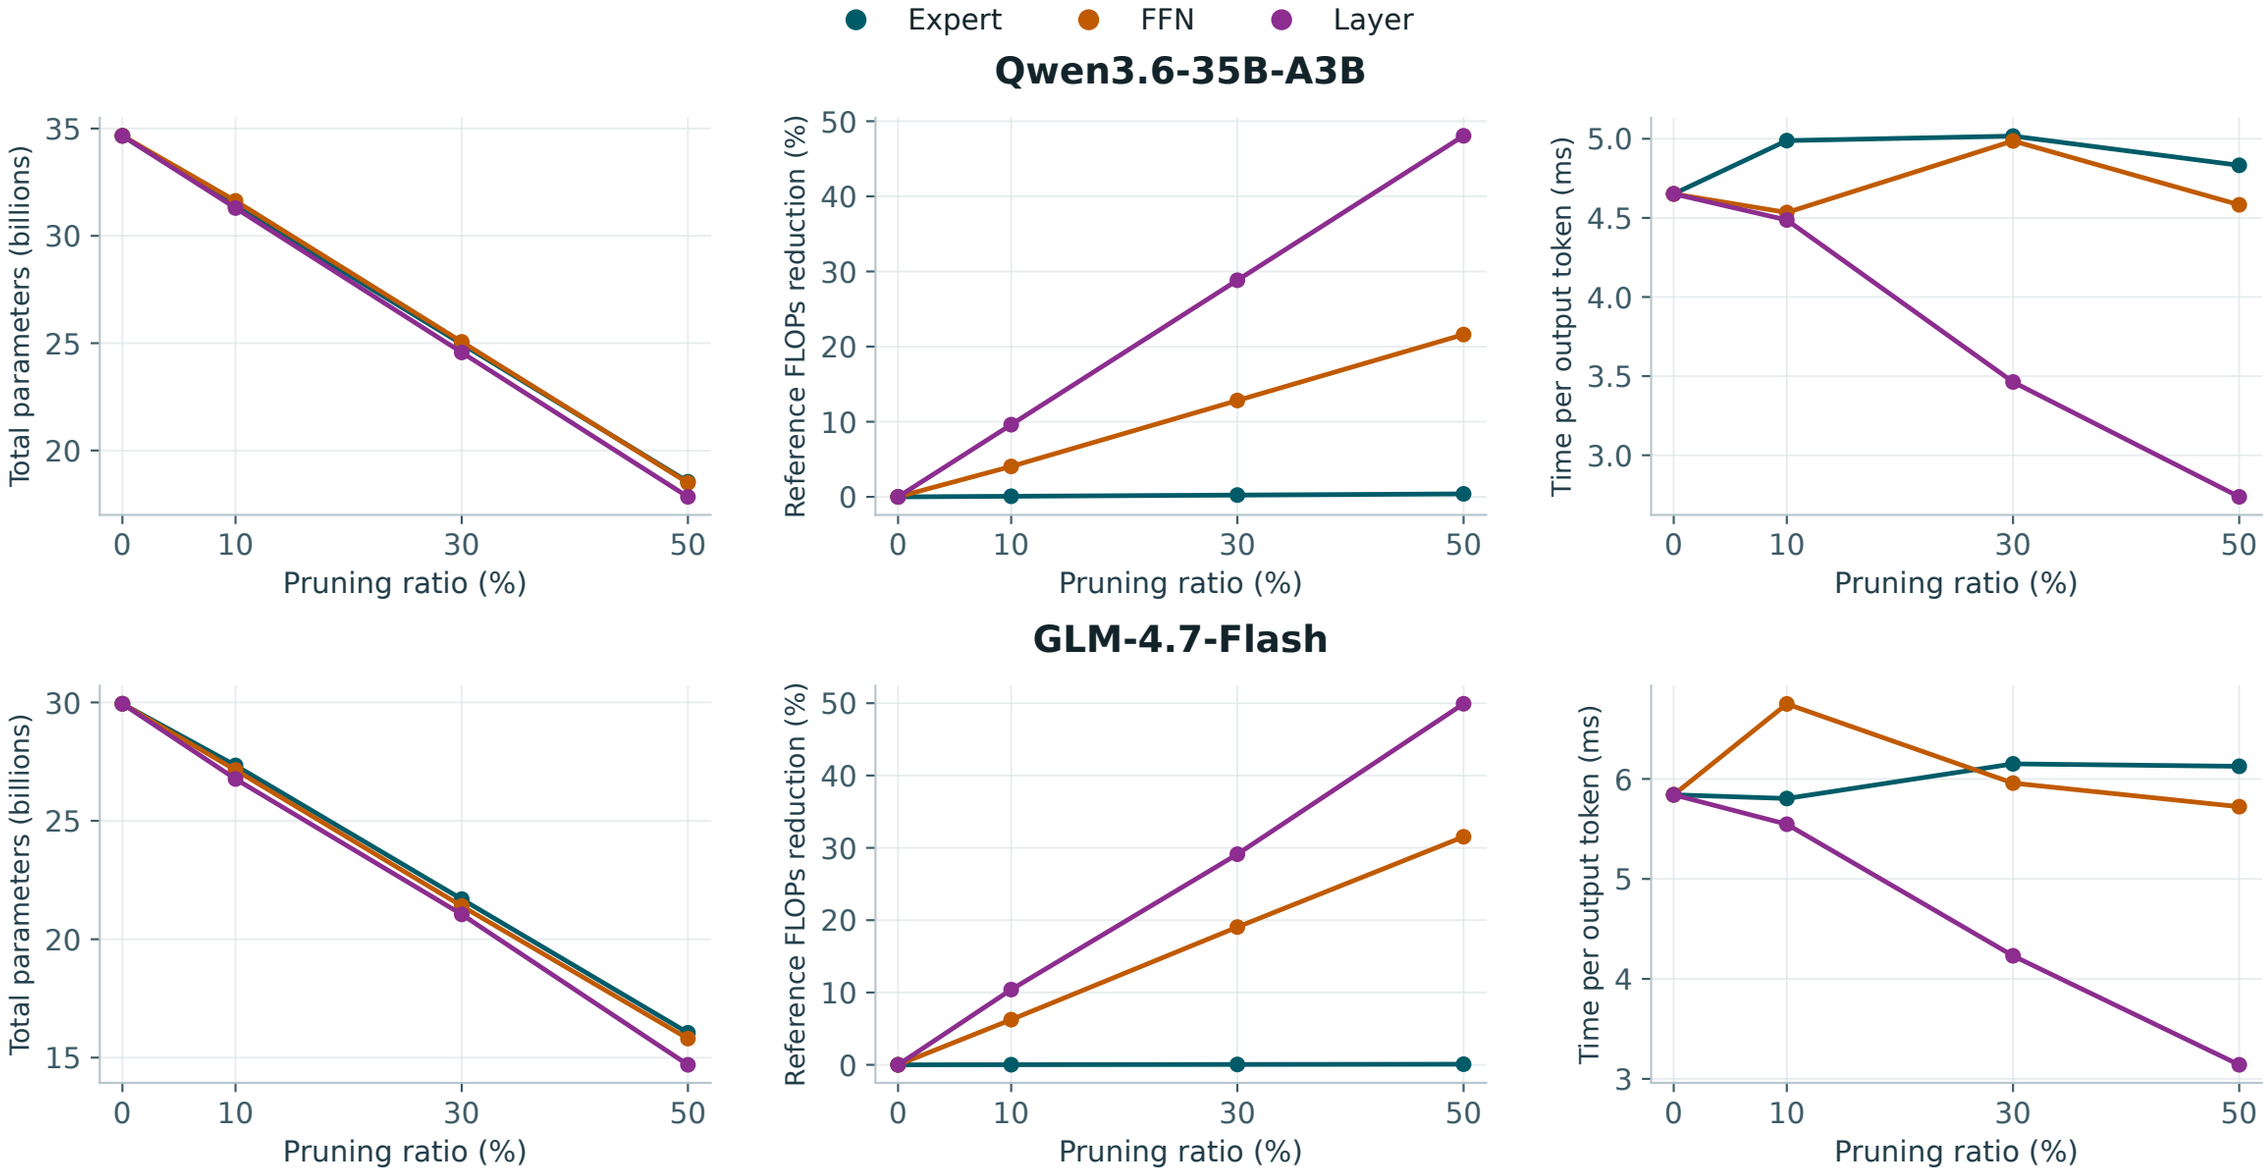


512 input tokens + exactly 128 generated tokens. Latency: median and interquartile range of five timed batches after two warmups.
Reference FLOPs are architecture-based estimates (one multiply-add = two FLOPs).