GLM-4.7-Flash | AIME 2026

Coverage and answer selection under the same cumulative output-token budget

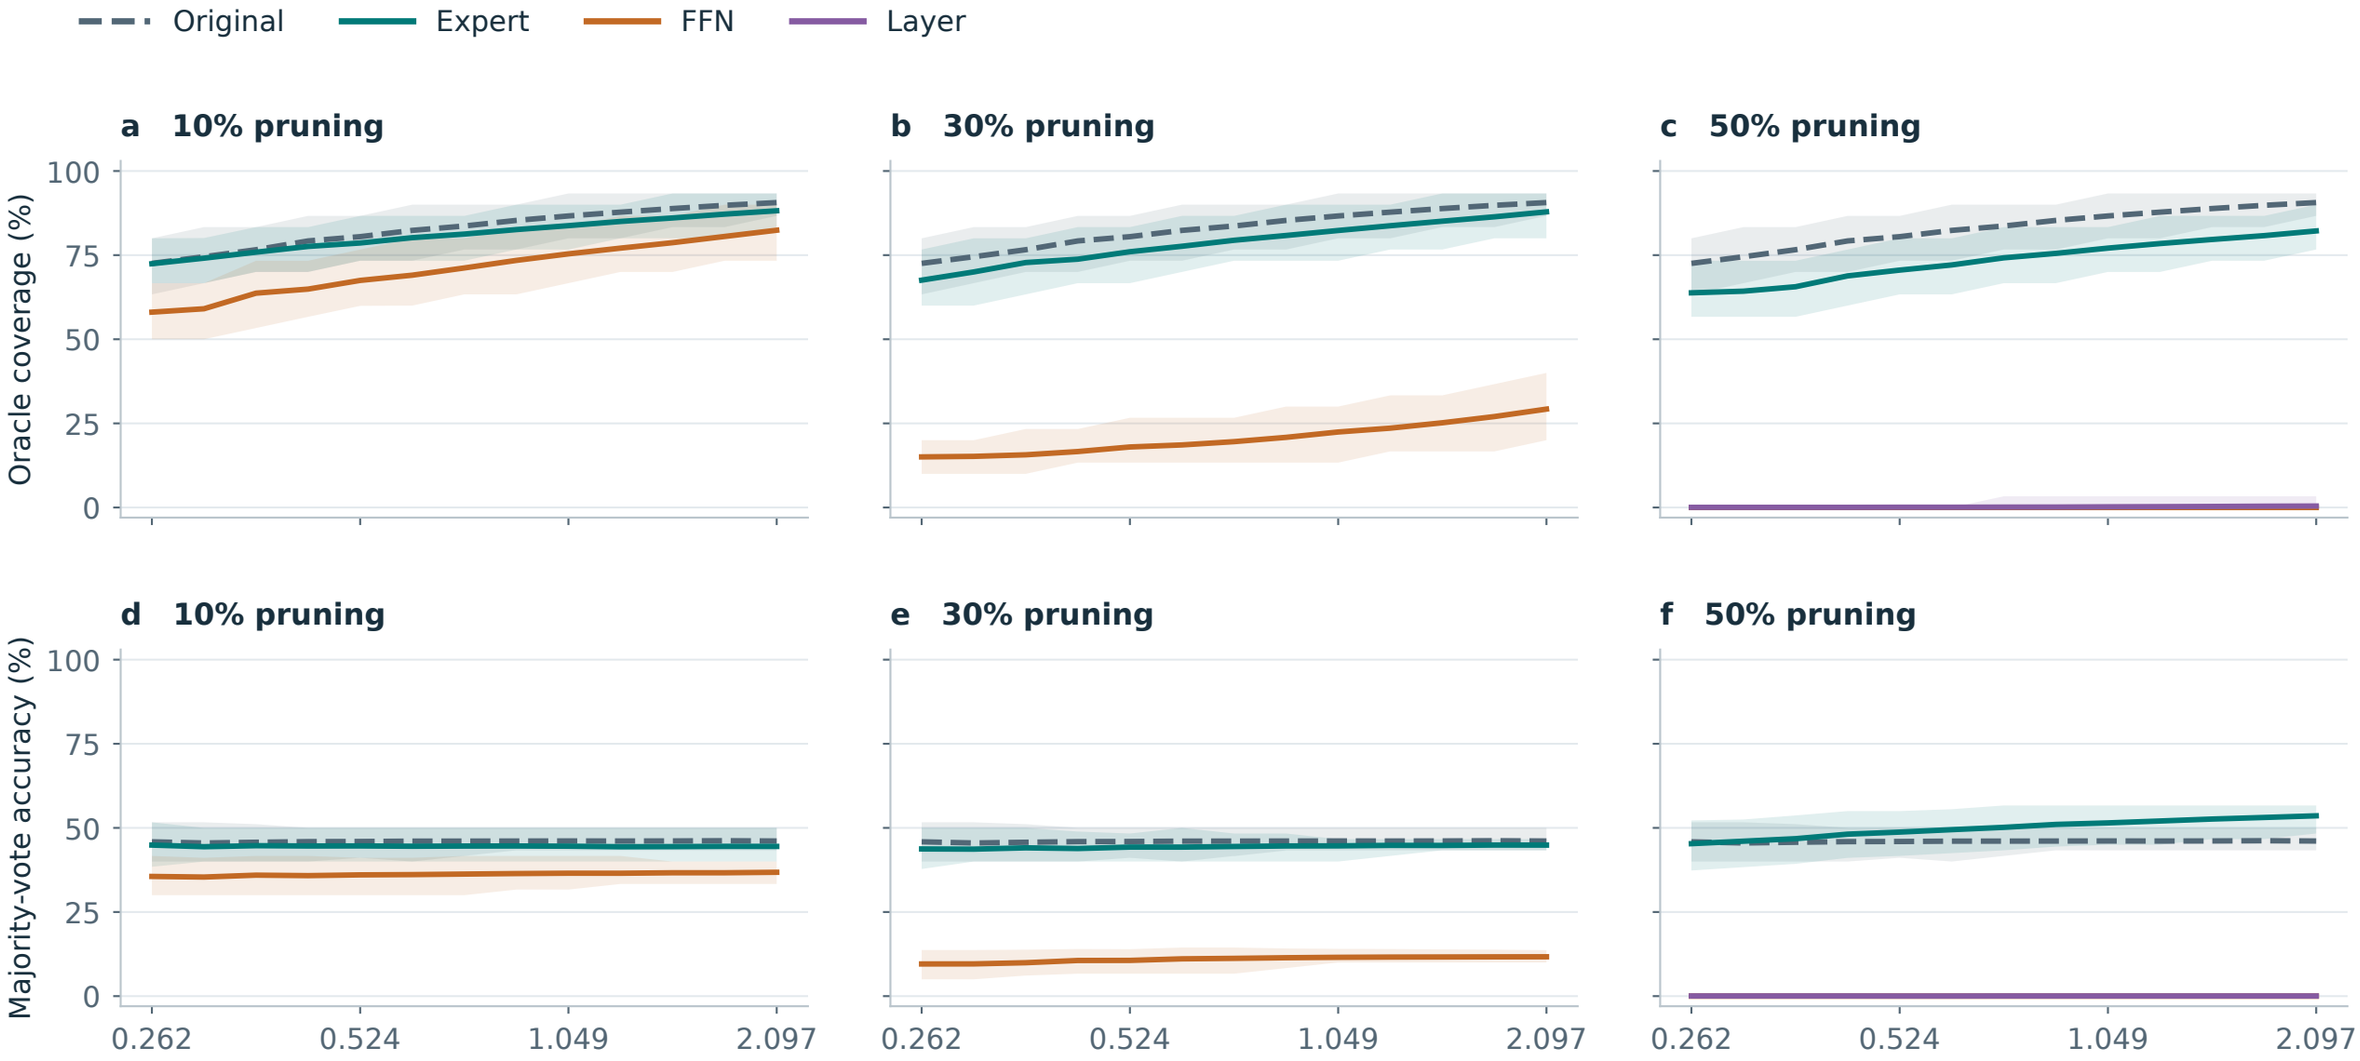

Lines: mean over 1,000 batch orderings. Shading: central 95% ordering range. Incorrect and truncated outputs incur cost.

Replay uses 32-candidate batches; each response has an 8,192-token limit. GLM Layer: 50% only; 10% and 30% unavailable.