Qwen3.6-35B-A3B | AIME 2026

Coverage and answer selection under the same cumulative output-token budget

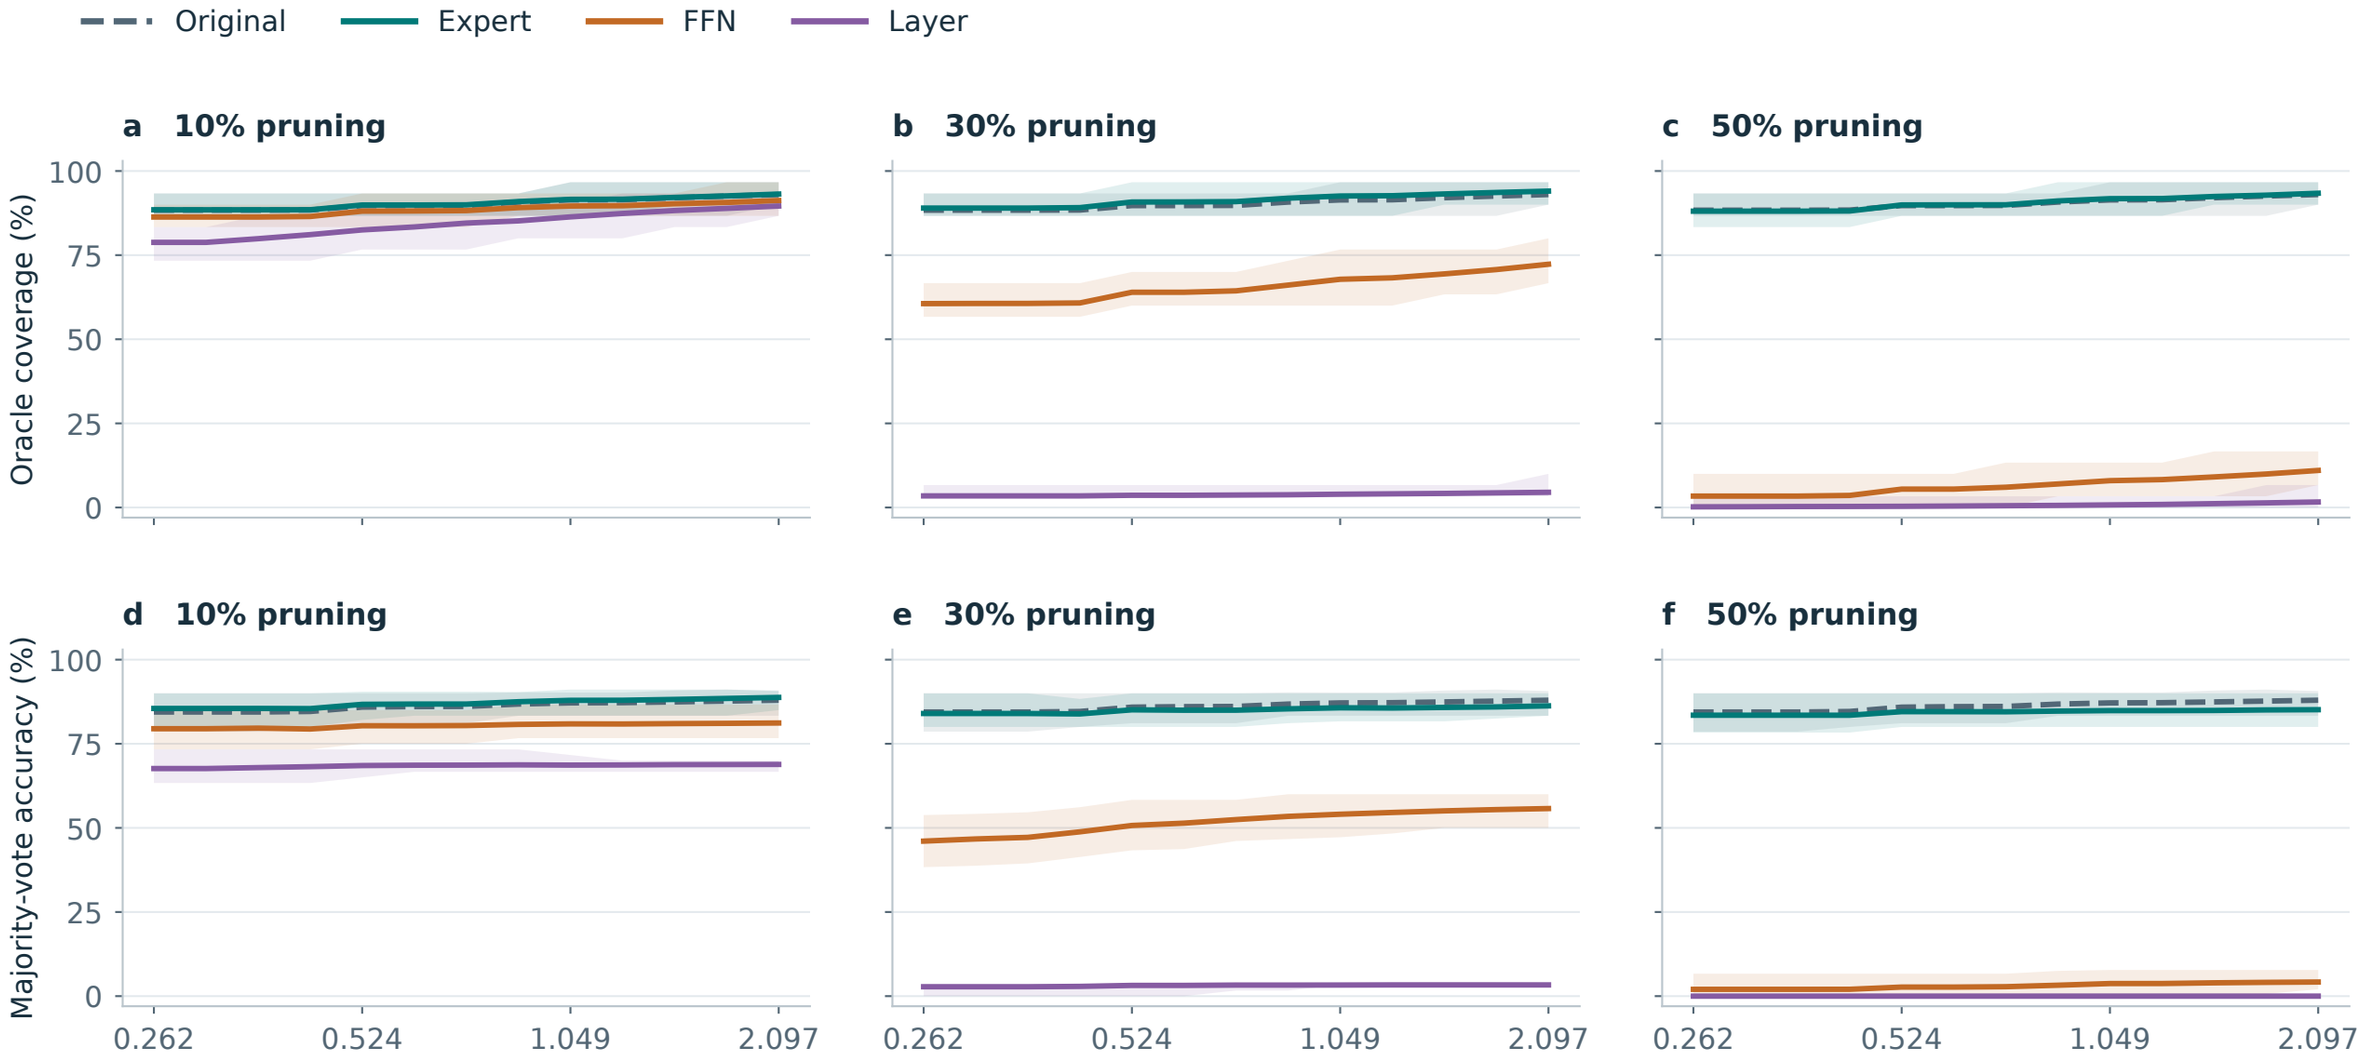

Lines: mean over 1,000 batch orderings. Shading: central 95% ordering range. Incorrect and truncated outputs incur cost.

Replay uses 32-candidate batches; each response has an 8,192-token limit.